



Supervised Learning: Regression, Classification & Time Series

Session 2

Bernardo Almada-Lobo, Daniel Pereira,
Pedro Amorim
Lisboa, 9 a 12 de outubro 2023

ACCREDITATIONS



MEMBERSHIPS



RANKINGS



Biography notes – Bernardo Almada-Lobo

Academic



Research



Consultancy

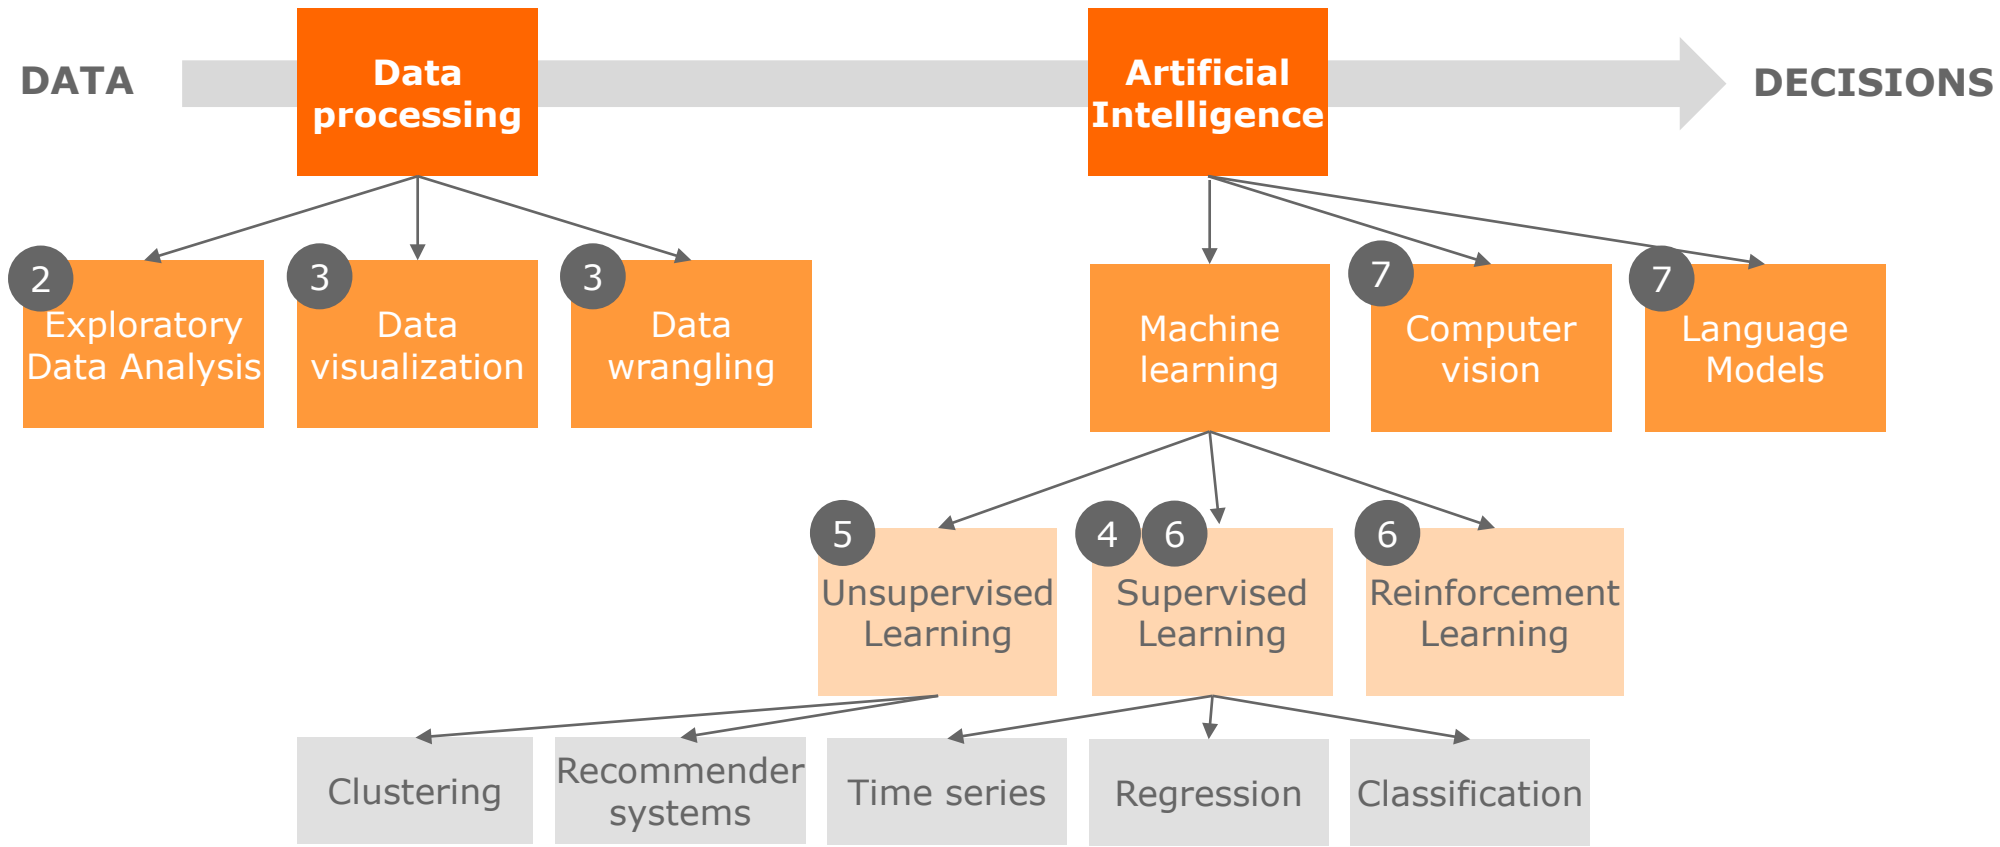


Education




Other






Algorithms can be classified based on the way they “learn” about data to make predictions


Supervised learning



Semi-supervised learning

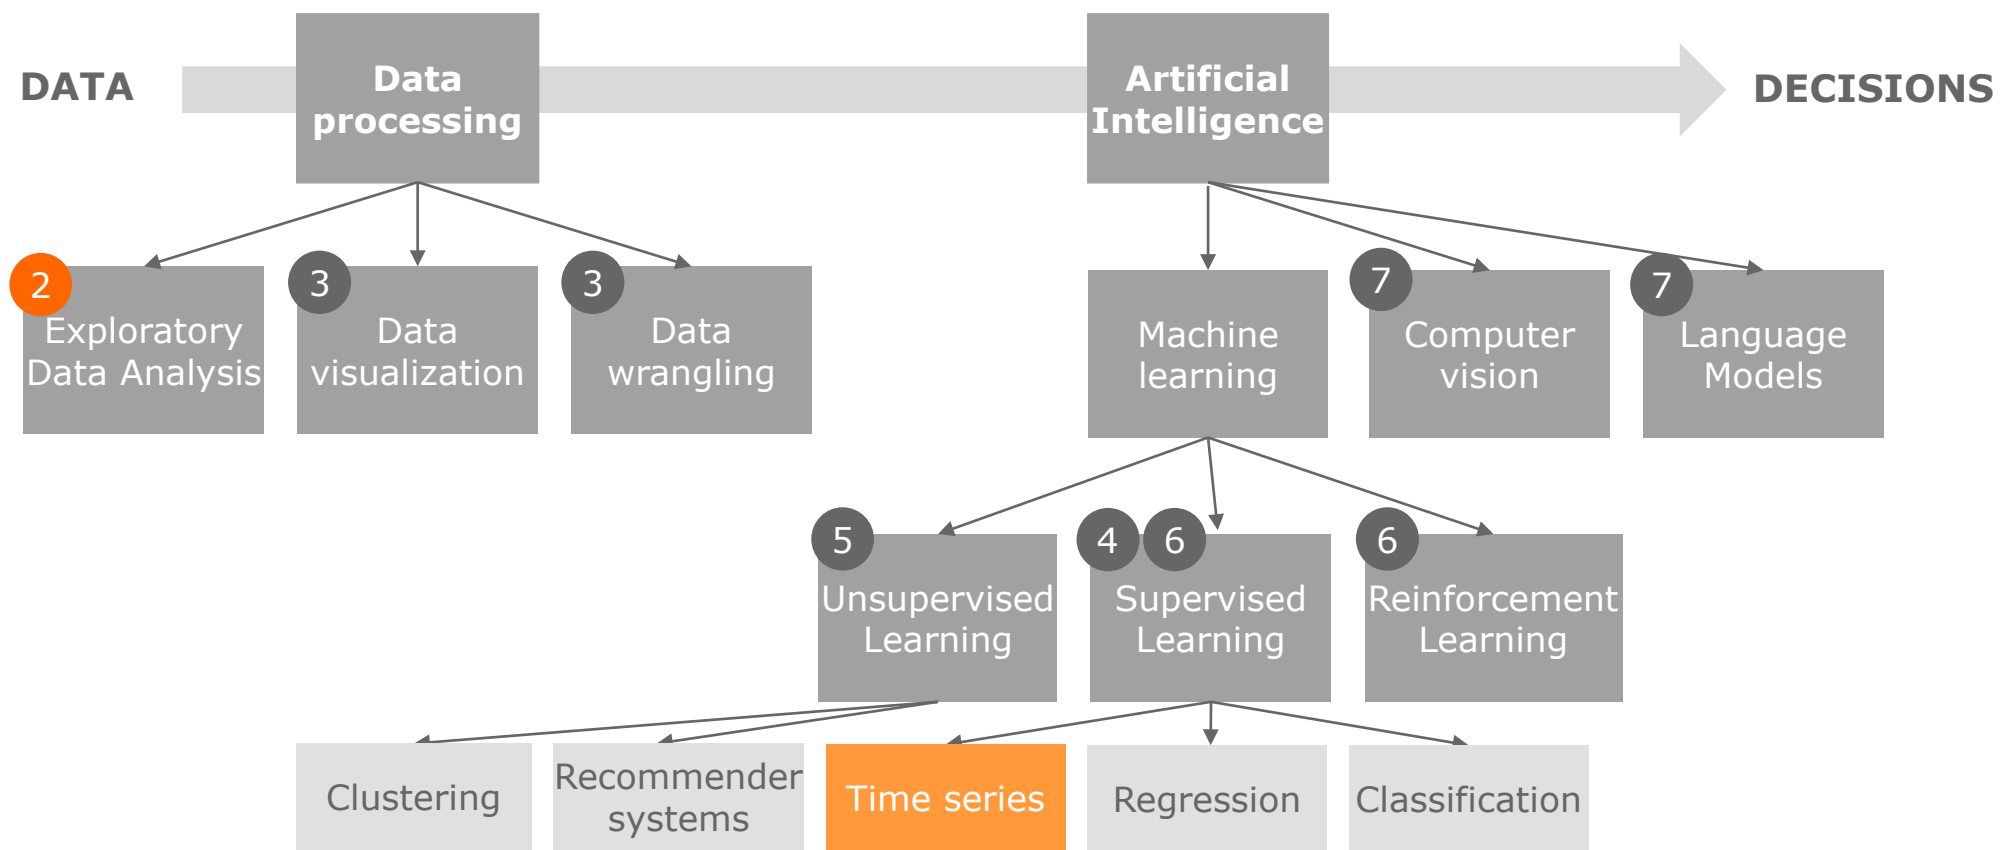


Unsupervised learning



Fundamentals

012



Google Cloud

8

The Role of Time Series

Time Series Methods:
Exponential Smoothing

Accuracy, Outliers and
Aggregation

Hybrid Methods

Box & Jenkins

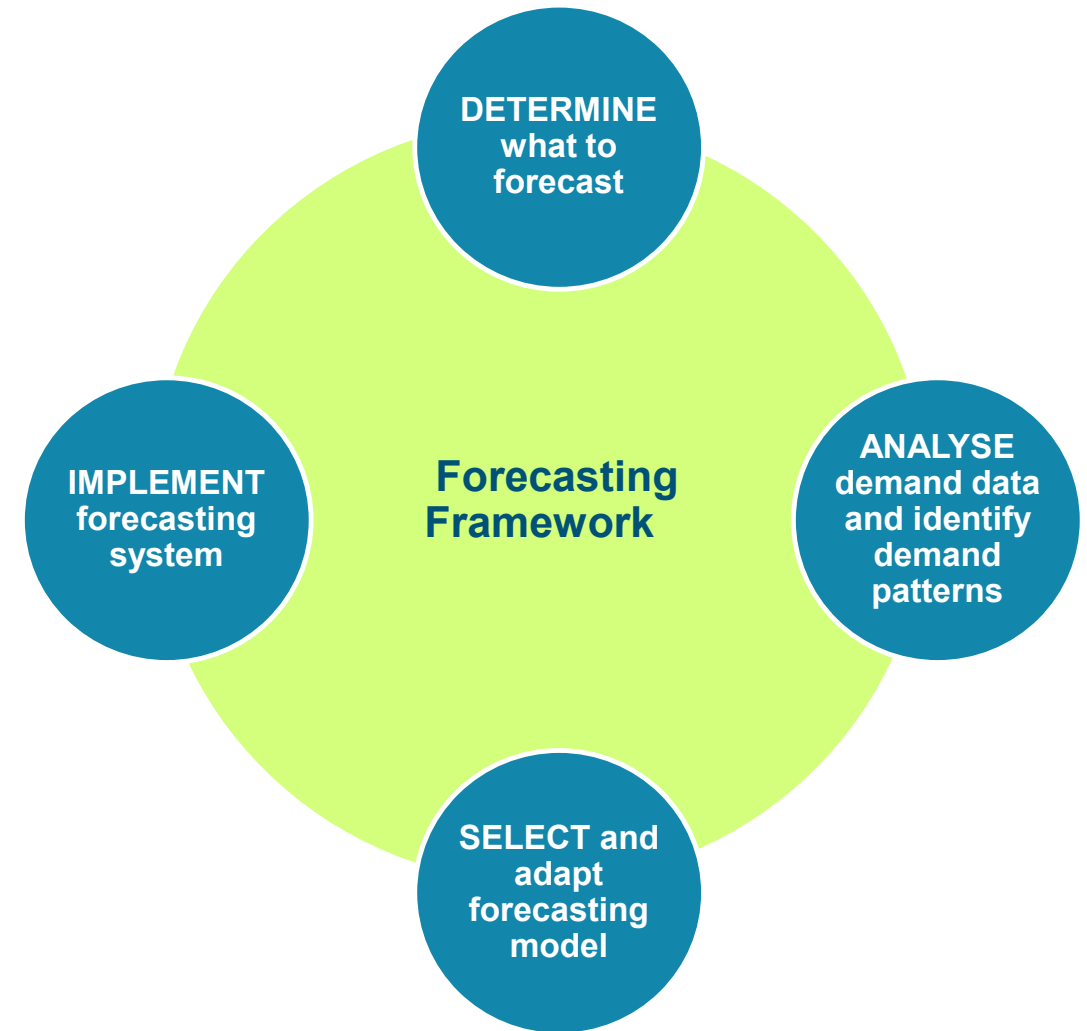


Why is Forecasting Important?

- Defining more precisely how ***UNCONTROLLABLE VARIABLES*** (relevant for decision process at stake) will behave in the future
- Lags in Decision Making: If we could always adjust instantaneously and costlessly to new conditions there would be no need for forecasts
- It is not an exact science; one must blend experience, judgment, and technical expertise

The six Steps in the Forecasting process


- **Step 1** Identify the goal of the forecast
- **Step 2** Establish a time horizon and the lag
- **Step 3** Select a forecasting technique
- **Step 4** Conduct the forecast (analyze data)
- **Step 5** Determine its accuracy
- **Step 6** Monitor the forecast



Forecasting Quantitative Methods


"It is far better to foresee even without certainty than not to foresee at all. "

Direct Extrapolation Methods



Casual Methods

Figure 1: Income versus years of education



Forecasting Quantitative Methods



The Role of Time Series

**Time Series Methods:
Exponential Smoothing**

Accuracy, Outliers and
Aggregation

Hybrid Methods


Box & Jenkins



Exercise

Consider the following time series:

t	Z_t
1	55
2	47
3	65
4	70
5	65
6	58
7	75
8	80
9	65
10	62
11	80
12	85
13	70
14	65
15	85
16	90




- What is the forecast for $t=17$?

A time series is composed of 3 main components: level, trend and seasonality

Decomposition models

Level

Sales




Time

Product expected average sales

Trend


Sales



Time

Time series evolution over time

Seasonality

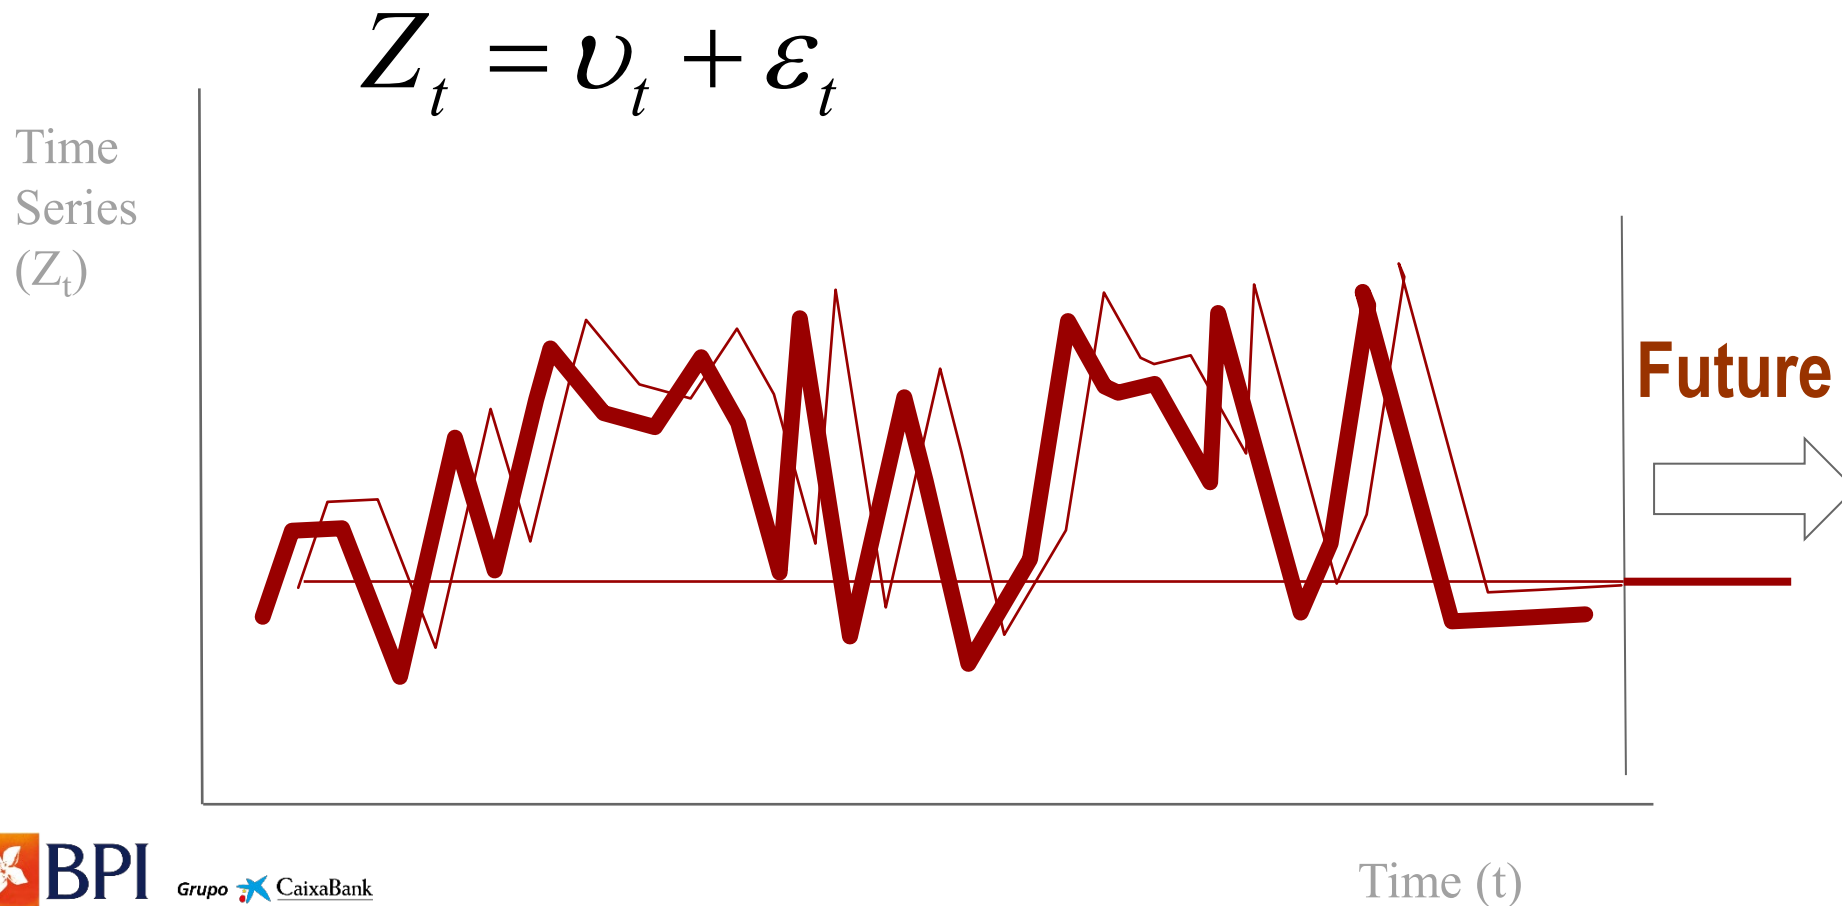


- **Sales patterns** on different periods – year, month, week

A decomposition model projects the impact of past **trend** and **seasonality** (additive or multiplicative) on a current sales **level** in order to forecast future sales

Exponential Smoothing Methods

Series Without Trend and Without Seasonality



Exponential Smoothing Methods

Local Stationary Time Series

Simple Moving Average Method

The level is estimated based on the last N observations

$$n_t = (Z_t + Z_{t-1} + \dots + Z_{t-N+1}) / N$$

The only parameter to fix is the number of terms of the moving average (N).


One could compute the value of N that minimizes the Mean Square Error of One Step Forecasts.

$$EQM = \frac{1}{T} \cdot \sum_{t=N+1}^{N+T} e_t^2 = \frac{1}{T} \cdot \sum_{t=N+1}^{N+T} [z_t - \hat{z}_{t-1}(1)]^2$$

Exponential smoothing shows valuable advantages when compared to moving average

Moving averaging models


Illustration



Pros/Cons


- + Simple method with rolling horizon
- Does not differentiate last N observations
- Ignores previous observations

Weighted moving average



- + Most recent observations are "heavier" (weight)
- Differentiates last N observations randomly
- Ignores previous observations

Exponential Smoothing



- + "Corrects", at each observation, the forecast value
- + Most recent observations are exponentially "heavier"
- + Uses all past observations
- + Only requires the last observation and the last forecast to be kept
- Implies parametrizing α

Exponential Smoothing Methods

Series Without Trend and Without Seasonality

Forecasts

$$\hat{Z}_t(1) = \hat{Z}_t(2) = \dots = \hat{Z}_t(k) = \dots = n_t = \hat{v}_t$$

$\hat{Z}_t(k)$: Forecast of Z_{t+k} made in instant t ,
after knowing the data Z_t .

n_t : Estimative of time series level
in instant t .

Simple Exponential Smoothing

$$n_t = \alpha \cdot Z_t + (1 - \alpha) \cdot n_{t-1} \quad (0 \leq \alpha \leq 1)$$

Since $\hat{Z}_t(1) = n_t$

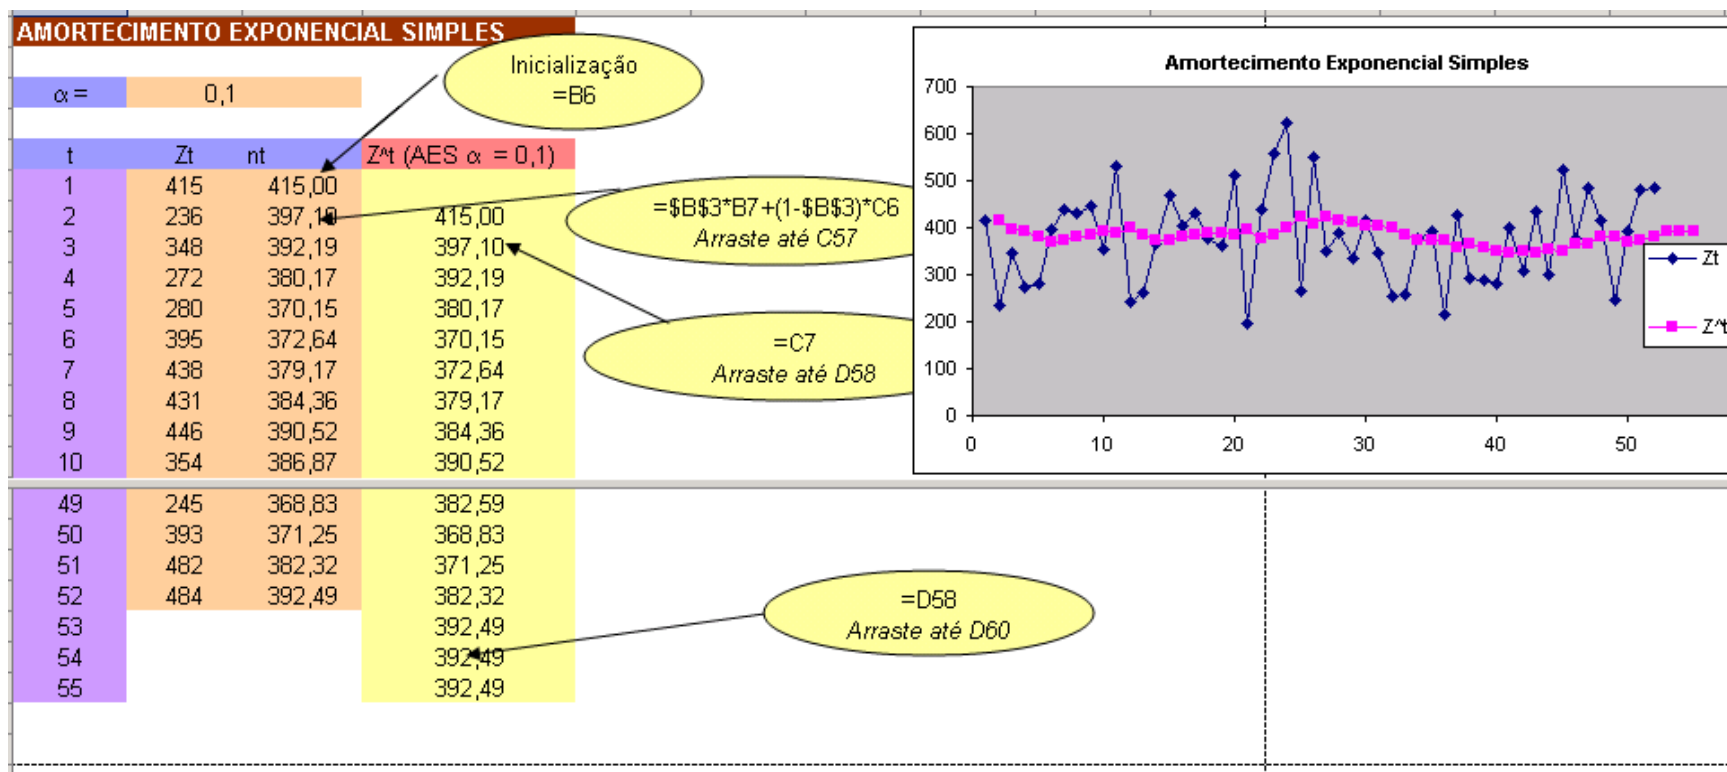
$$\begin{aligned} \hat{Z}_t(1) &= \alpha \cdot Z_t + (1 - \alpha) \cdot \hat{Z}_{t-1}(1) \\ &= \hat{Z}_{t-1}(1) + \alpha \cdot [Z_t - \hat{Z}_{t-1}(1)] \\ &= \hat{Z}_{t-1}(1) + \alpha \cdot e_t \quad (0 \leq \alpha \leq 1) \end{aligned}$$

α : Smoothing rate

e_t : Forecast error in instant t

Exponential Smoothing Methods

Simple Exponential Smoothing





Smoothing factors determine how reactive forecast model is towards demand variations

Exponential smoothing models

— Sales — Forecast



Example



Exponential Smoothing Methods

Local stationary Time Series: Tracking Signal

Adaptive Exponential Smoothing de Trigg & Leach

$$TS_t = \frac{EA_t}{EAA_t}$$

$$EA_t = \beta \cdot e_t + (1 - \beta) \cdot EA_{t-1} \quad (\text{Smoothed Error})$$

$$EAA_t = \beta \cdot |e_t| + (1 - \beta) \cdot EAA_{t-1} \quad (\text{Absolute Smoothed Error})$$

$$e_t = Z_t - \hat{Z}_{t-1} \quad (1) \quad (\text{Forecast Error})$$

β : *smoothing rate* (constant) of forecast error and absolute forecast error

$$\alpha_t = |TS_t|$$

Naive Forecasts (equal to $\alpha=1$)


- ✓ A naive forecast for any period equals the previous period's actual value
- ✓ Low cost, easy to prepare, easy to understand, but less accurate forecasts
- ✓ Can be applied to seasonal or trend data

Examples:

- If last week's demand was 50 units, the naive forecast for the coming week is 50 units.
- If seasonal pattern exists, the naive forecast for *next January* would equal the actual demand for *January of this year*.

Exponential Smoothing Methods

Series With Trend and Without Seasonality



Exponential Smoothing Methods

Series With Trend and Without Seasonality

Forecasts

$$\hat{Z}_t(k) = n_t + b_t \cdot k$$

$\hat{Z}_t(k)$: Forecast of Z_{t+k} made in instant t , after knowing the data Z_t

$n_t = \hat{v}_t$: Estimate of the level of the series in instant t .

$b_t = \hat{\tau}_t$: Estimate of the trend of the series

Linear Holt Exponential Smoothing

$$n_t = \alpha \cdot Z_t + (1 - \alpha) \cdot (n_{t-1} + b_{t-1}) \quad 0 \leq \alpha \leq 1$$

$$b_t = \beta \cdot (n_t - n_{t-1}) + (1 - \beta) \cdot b_{t-1} \quad 0 \leq \beta \leq 1$$

$$\hat{Z}_t(k) = n_t + b_t \cdot k$$

Initialization:

$$b_1 = 0$$

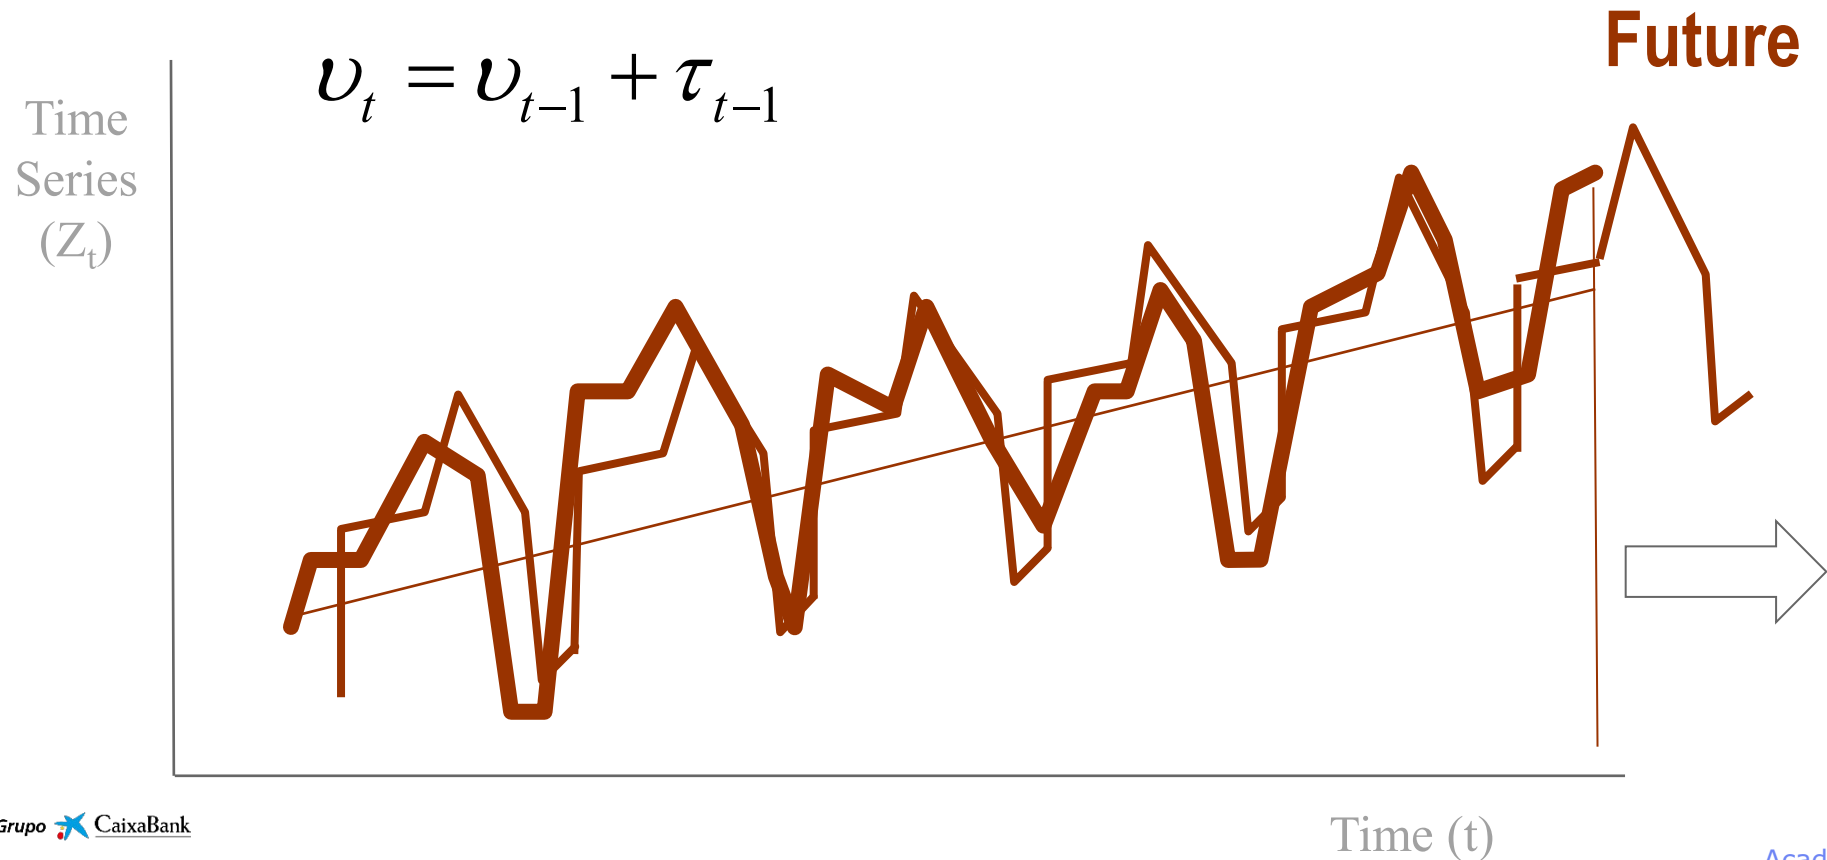
$$n_1 = Z_1$$

Exponential Smoothing Methods

Series With Trend and Seasonality

$$Z_t = v_t + \phi_t + \varepsilon_t$$

$$v_t = v_{t-1} + \tau_{t-1}$$




Besides level, ES models may consider trend and seasonality as well


Exponential smoothing models

— Sales — Forecast


Simple Exponential Smoothing
(Simple ES)



Holt's Method
(Trend ES)



Holt Winters' Method
(Seasonal ES)



Exponential Smoothing Methods

HoltWinters: Series With Trend and Seasonality

Forecasts

$$\hat{Z}_t(k) = n_t + b_t \cdot k + f_{t+k-s}$$

$$\hat{Z}_t(k)$$

Forecast of Z_{t+k} made in instant t ,
after knowing the data Z_t

$$n_t = \hat{v}_t$$

Estimate level of series in instant t .

$$b_t = \hat{\tau}_t$$

Estimate of the trend of the series.

$$f_{t+k-s} = \hat{\phi}_{t+k}$$

Estimate of the seasonal component
for instant $t+k$.

Holt-Winters Method – Additive Model

$$n_t = \alpha \cdot (Z_t - f_{t-s}) + (1 - \alpha) \cdot (n_{t-1} + b_{t-1}) \quad 0 \leq \alpha \leq 1$$

$$b_t = \beta \cdot (n_t - n_{t-1}) + (1 - \beta) \cdot b_{t-1} \quad 0 \leq \beta \leq 1$$

$$f_t = \gamma \cdot (Z_t - n_t) + (1 - \gamma) \cdot f_{t-s} \quad 0 \leq \gamma \leq 1$$

$$\hat{Z}_t(k) = n_t + b_t \cdot k + f_{t+k-s}, \text{ for } k = 1, 2, \dots, s$$

$$\hat{Z}_t(k) = n_t + b_t \cdot k + f_{t+k-2s}, \text{ for } k = s+1, s+2, \dots, 2s$$

Initialization:

$$i) \quad n_s = Z^* = \frac{1}{s} \cdot \sum_{t=1}^s Z_t$$

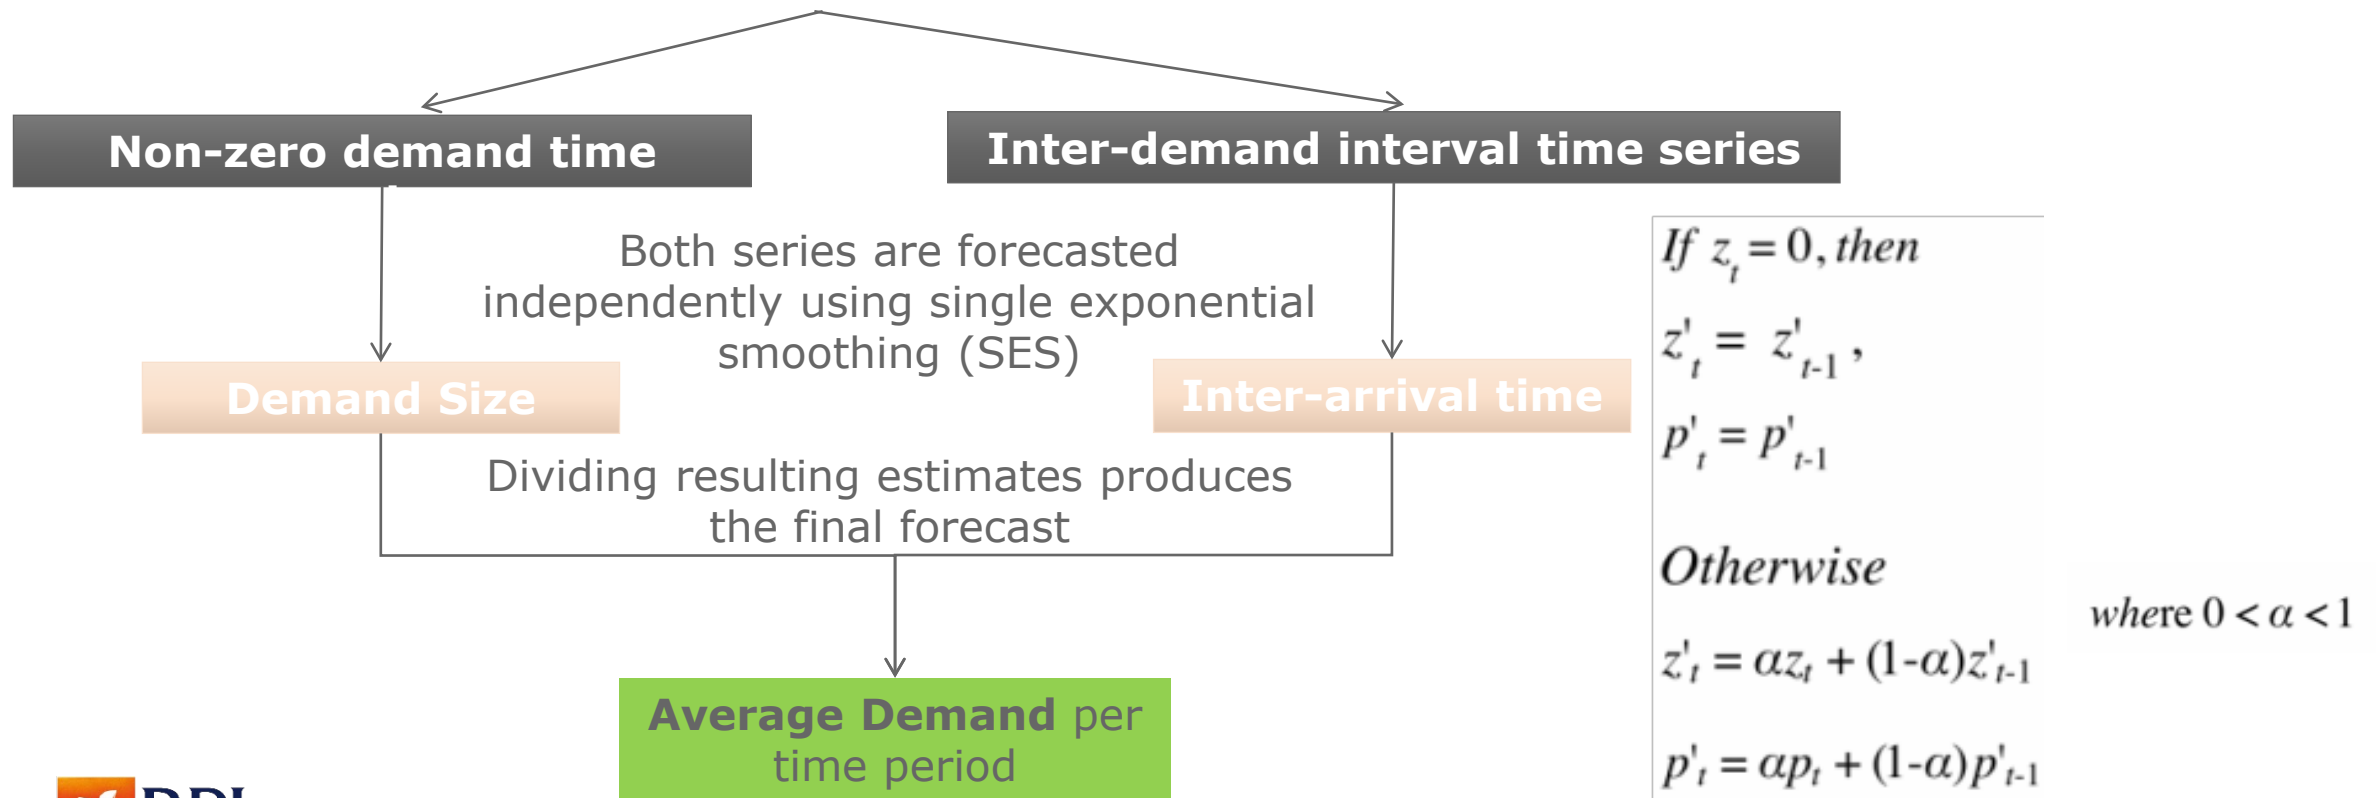
$$ii) \quad b_s = 0$$

$$iii) \quad f_j = Z_j - Z^* \quad (j = 1, \dots, s)$$

Croston Method for Intermittent Demand

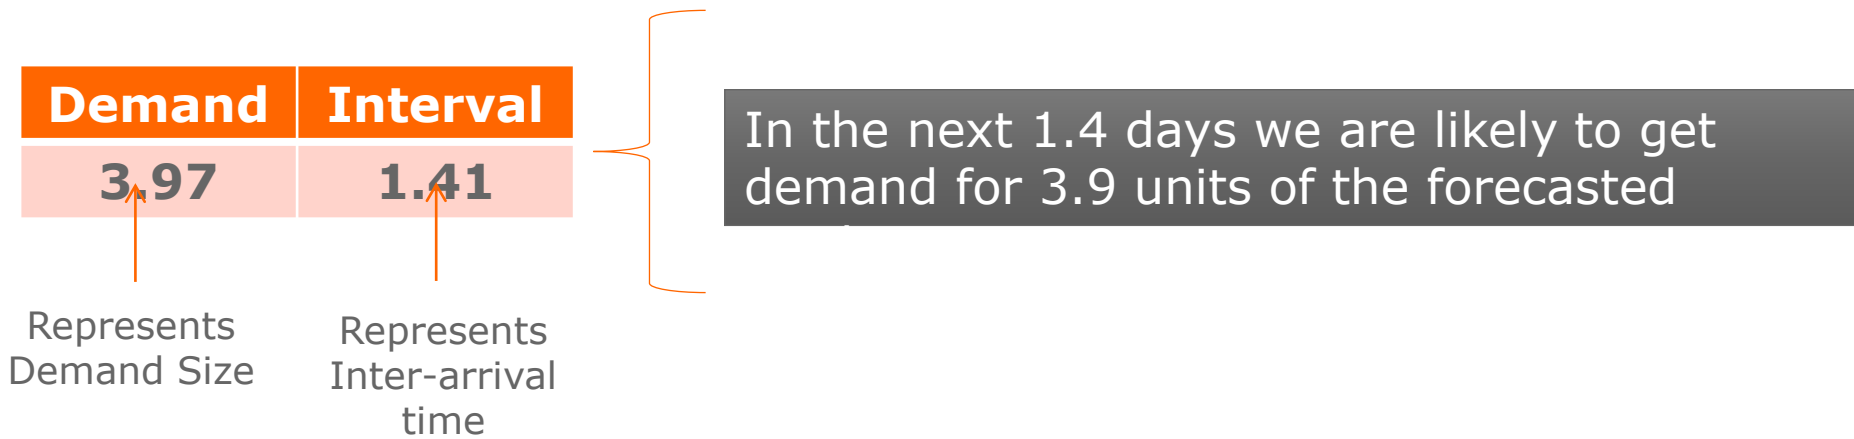
When a product experiences several periods of zero demand. Often demand is small, and sometimes highly variable in size.

Separates intermittent data into two components




Interpreting Croston Results

Sample outputs from Croston



Assuming we are forecasting for next 5 days, below will be the output for average demand for these 5 days.



The Role of Time Series

Time Series Methods:
Exponential Smoothing

**Accuracy, Outliers and
Aggregation**

Hybrid Methods

Box & Jenkins



Forecast's error impacts the entire planning activities and lead to future uncertainty


Why does this topic matter?

— Sales — Forecast

Forecast with **positive** BIAS (overshooting)

What's the expected impact of such forecast's behavior?

Sales




Time

Forecast with **negative** BIAS (undershooting)

What's the expected impact of such forecast's behavior?

Sales



Time

Forecast errors impacts companies in different ways

Example of a “fresh” food retailer



Stockouts



Lack of Visibility



Depreciated sales



Shrinkage

Sales Underestimation

- **Lost sales** potential
- **Reduced** customer loyalty and **satisfaction**

Sales Overestimation

- Unnecessary **stock** level **increase**
- **Increase** in capital, transportation and warehousing **costs** regarding **stock**
- **Higher risk of expiring** validity (shrinkage)

An ABC/XYZ analysis allows to differentiate products based on their relevance in sales and difficulty to forecast

Which products should be forecasted by an analytical model?

		X	Y	Z
A				
B				
C				

20% products

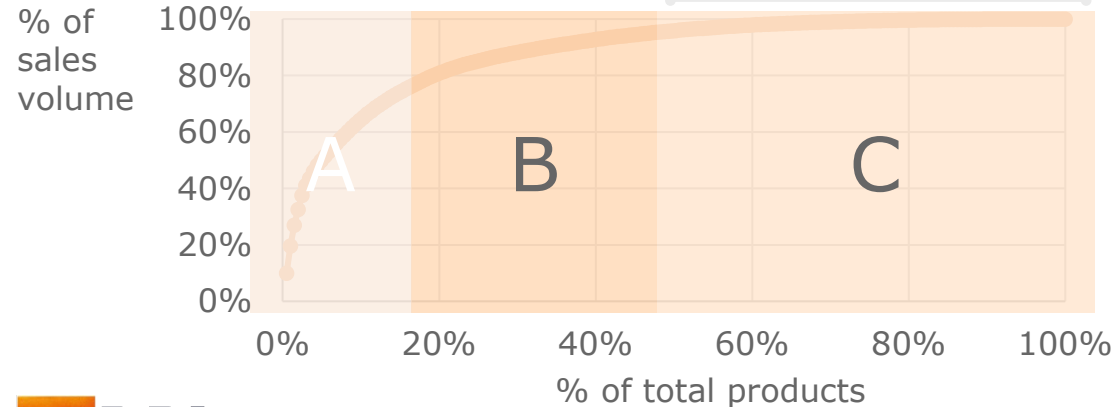
80% sales

20-50% products

15% sales


50-100% products

5% sales



Difficulty to forecast

	X	Y	Z
A			
B			
C			




$$COV = \frac{\text{Standard Deviation of Demand}}{\text{Mean of Demand}}$$

X

Relatively stable, small fluctuations or with a clear pattern

High forecast accuracy



Y

More significant fluctuations


Medium forecast accuracy¹



Z

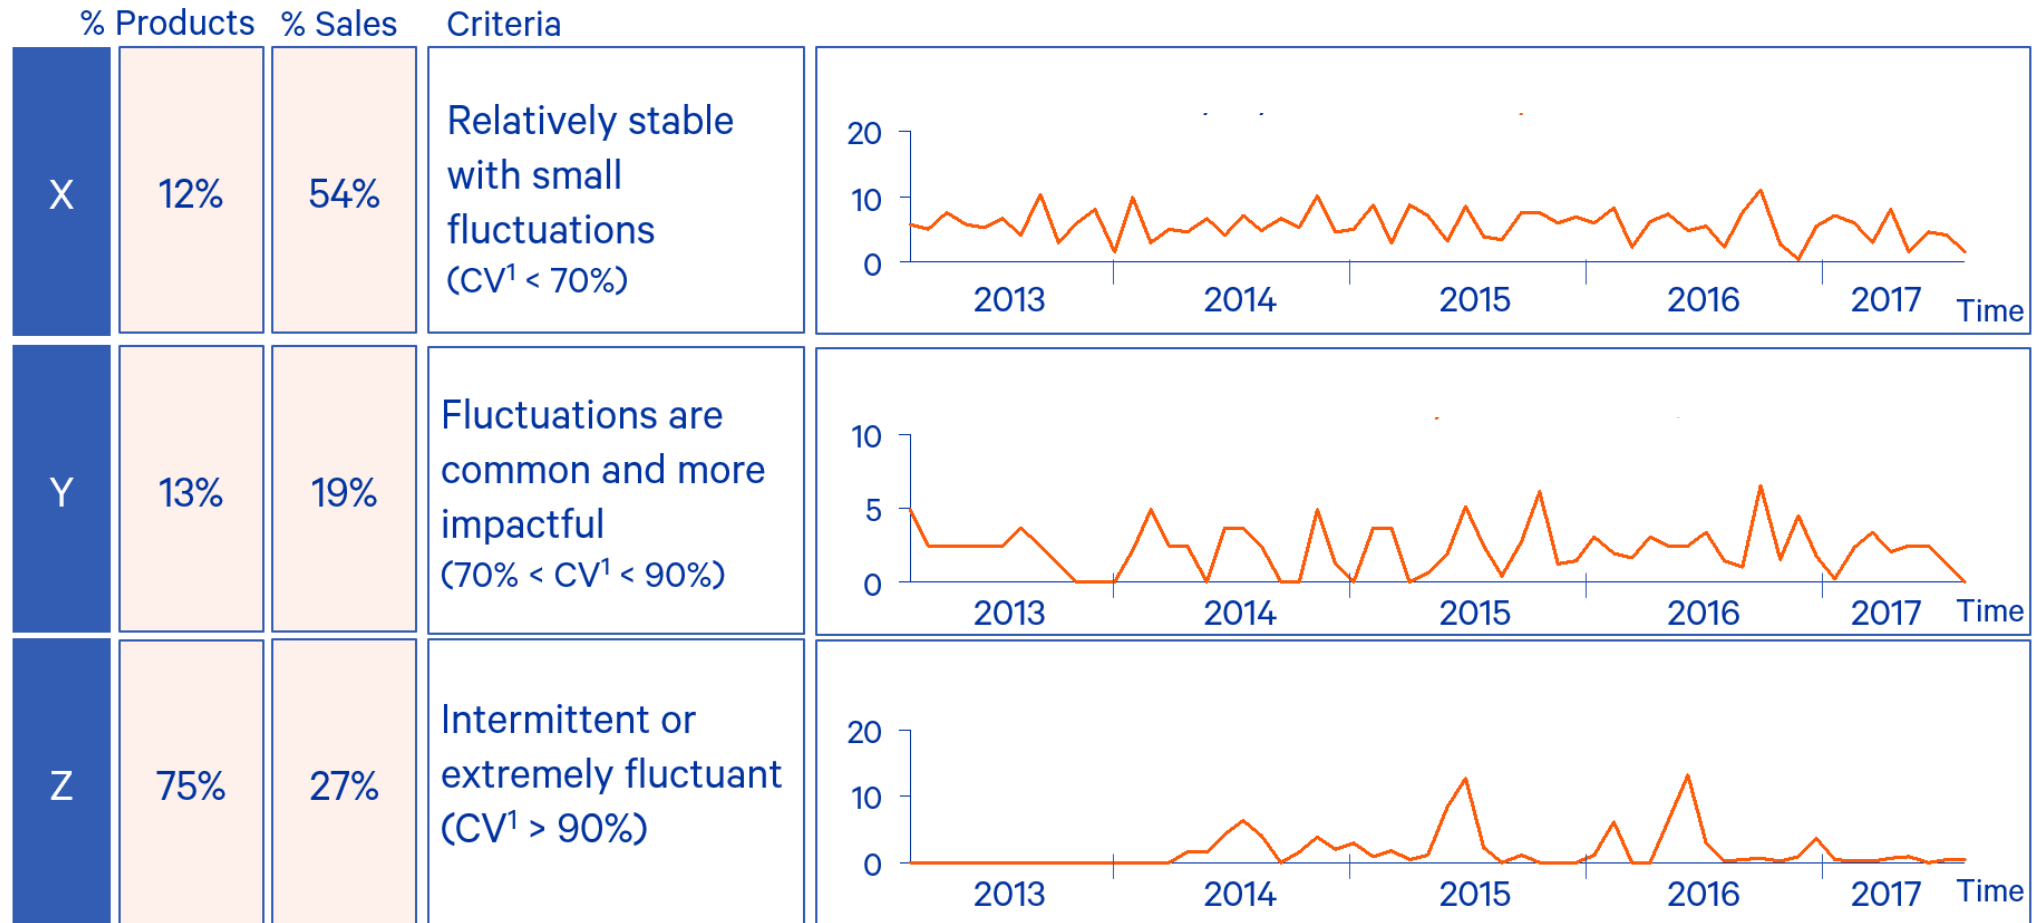
Very intense and intermittent fluctuations, no clear pattern

Low forecast accuracy



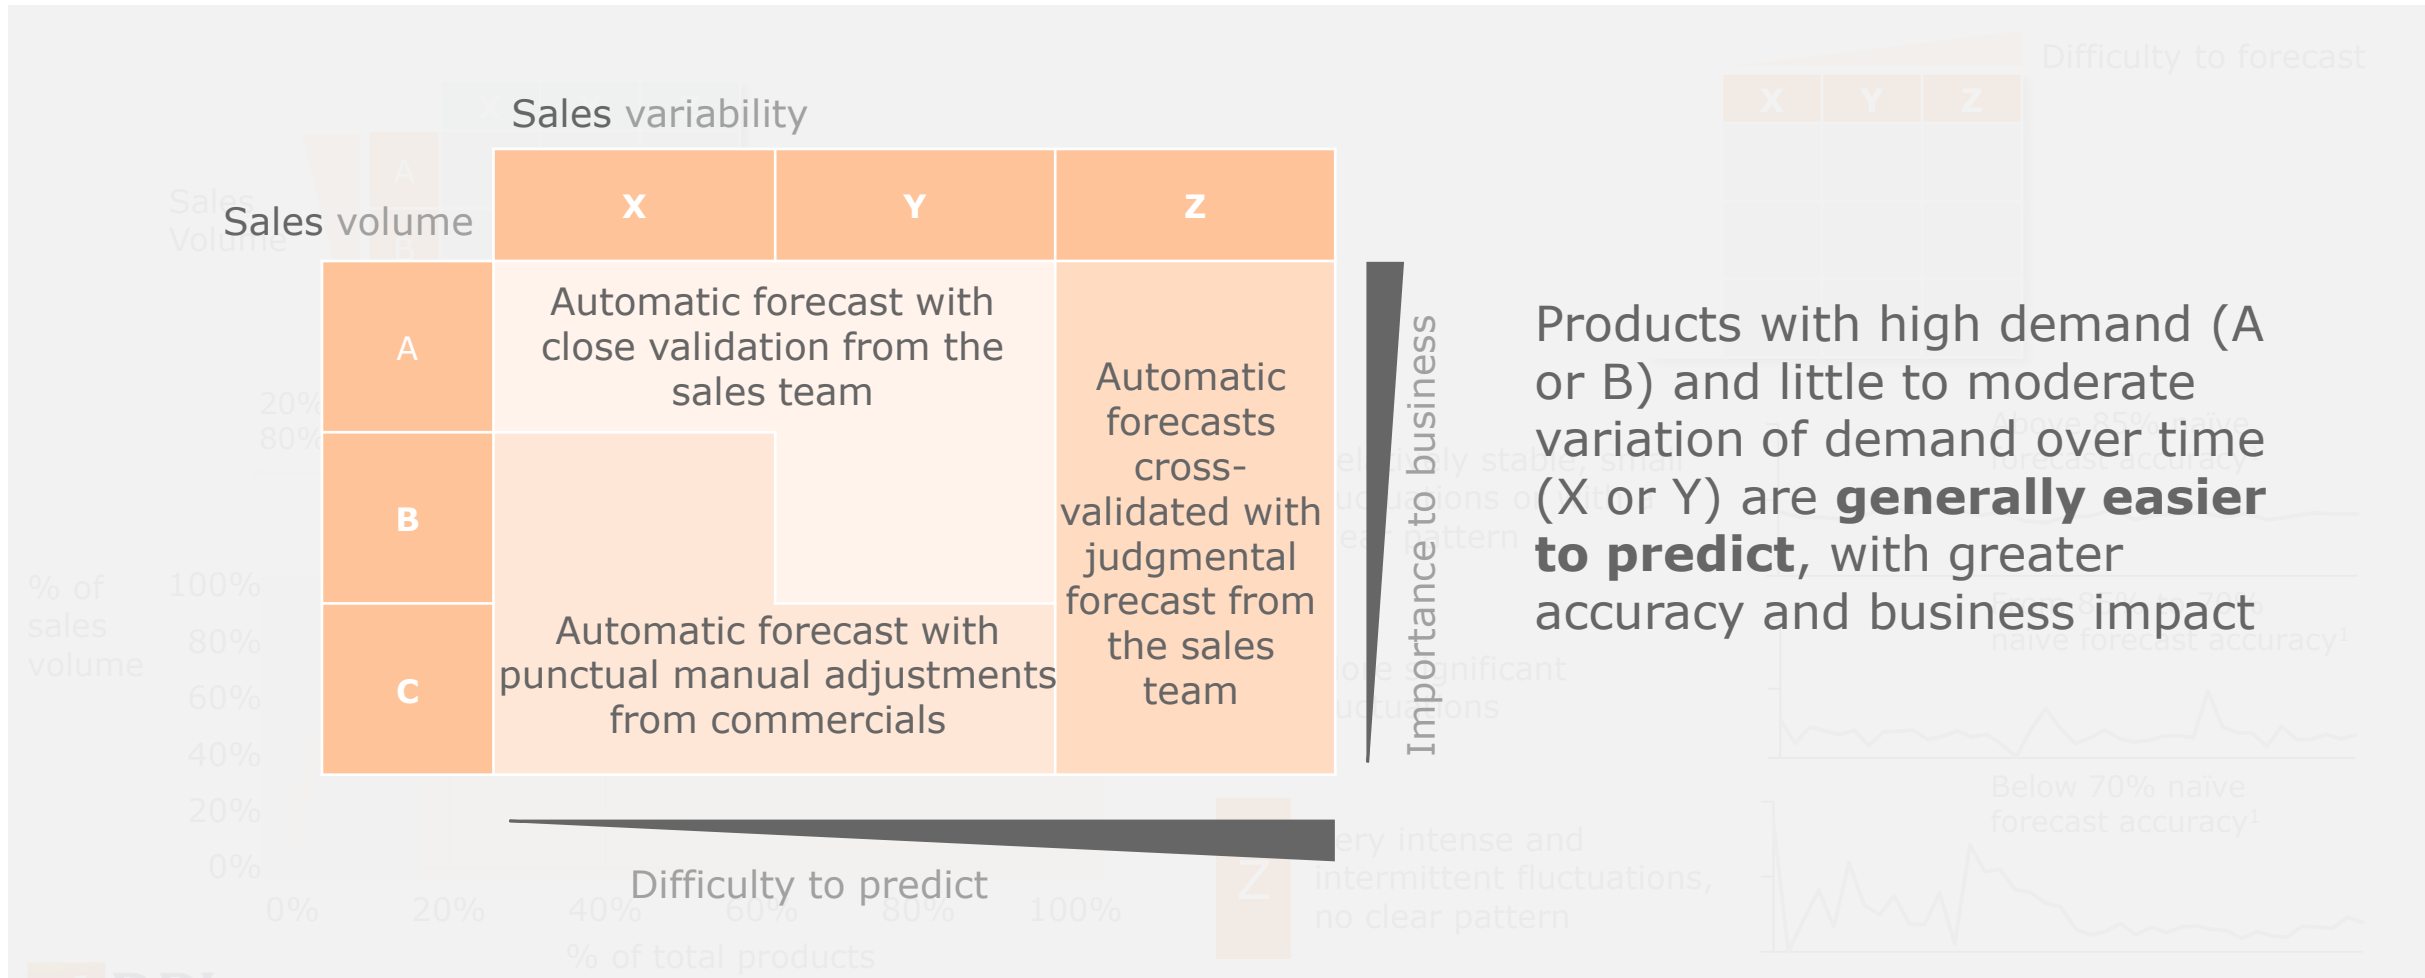
An XYZ analysis differentiates products over their variability and thus the difficulty of forecast generation

Example#1



An ABC/XYZ analysis allows to differentiate products based on their relevance in sales and difficulty to forecast

Forecast evaluation



Exercise ABC/XYZ: Italian Retailer

Open Sales Data.xls



Commercial inputs are increasingly important for Y and Z products

Example#2

ABC/XYZ analysis

%analyzed SKUs 60,3%

%in total sales 96,6%




¹ accuracy measured at SKU x month level, system accuracy in brackets

Evaluating forecast's performance is crucial in order to fine-tune the predictive models

Forecast evaluation


High bias,
High variability




High bias,
Low variability



Low bias,
High variability



Low bias,
Low variability



Bias: systematic overestimation (or underestimation)

Variability: random error (unpredictable)

There is a variety of metrics that can be used to evaluate forecast's precision, both relative and absolute

Absolute Measures

$$e_t = \hat{Z}_t - Z_t$$

$$ME = \frac{\sum_{t=1}^n e_t}{n}$$

Mean Error

$$MAE = \frac{\sum_{t=1}^n |e_t|}{n}$$

Mean Absolute Error

$$MSE = \frac{\sum_{t=1}^n e_t^2}{n}$$

Mean Squared Error

Relative Measures

$$PE_t = \left(\frac{\hat{Z}_t - Z_t}{Z_t} \right) \times 100$$

Percentual Error

$$BIAS(MPE) = \frac{\sum_{t=1}^n PE_t}{n}$$

Mean Percentual Error

$$MAPE = \frac{\sum_{t=1}^n |PE_t|}{n}$$

Mean Absolute Percentual Error

MAPE and **BIAS** are the most common metrics to evaluate forecast's precision. These metrics are weighted by the sales in order to guarantee that forecasting errors in high seller products contribute more significantly for the global error

Metrics may be weighted given the importance of each product

Formula

Example

Simple Average

$$BIAS(EPM) = \frac{\sum_i^N BIAS_i}{N}$$

$$MAPE = \frac{\sum_i^N MAPE_i}{N}$$

	<i>i=1</i>	<i>i=2</i>	<i>i=3 (N)</i>	
	A	B	C	
Vendas	10 000	500	20	MAPE=60%
MAPE:	10%	70%	100%	

Weighted Average

$$BIAS(EPM) = \frac{\sum_i^N (\hat{Z}_i - Z_i)}{\sum_i^N Z_i}$$

$$MAPE = \frac{\sum_i^N |\hat{Z}_i - Z_i|}{\sum_i^N Z_i}$$

	A	B	C	
Vendas:	10 000	500	20	MAPE=13%
MAPE:	10%	70%	100%	

Note: Z_t – Real sales at period t ; \hat{Z}_t – Sales forecast, at $t-1$, for period t

MAPE and BIAS provide distinct, rich information at different aggregation levels

Forecast evaluation

— Sales — Forecast





MAPE - Mean absolute percentual error evaluates the mean deviation at a detailed level




BIAS - Forecast bias identifies systematic deviations, either underestimation or overestimation

Aggregated level
(e.g. group of products)




Detailed level
(e.g. product)


MAPE = 22.4%
BIAS = + 16.7%



MAPE = 14.5%
BIAS = - 10.5%



MAPE = 16.1%
BIAS = + 11.8%



Dashboard de Análise de *Forecasts* Regular

Relatório Semanal

dc_desc

0010 ALIMENTAR

0011 PEIXARIA&TALHO

0012 F&L C&Q PAD
TAWAY

0021 CASA

0023 BAZAR

0041 NUTRIÇÃO
SAUDÁVEL

eow

4 de May de 2020

Evolução semanal do MAPE

45.2% 42.5%



● MAPE Regular Aprovado: Retek ● MAPE Regular RDF: Sistema

Variação do MAPE aprovado (pp)

Variação face à semana anterior


● -4.13

Variação face às últimas 6 semanas

● -4.07

Evolução semanal do BIAS

2.7% 2.0%



● BIAS Regular Aprovado: Retek ● BIAS Regular RDF: Sistema

Variação do BIAS aprovado (pp)

Variação face à semana anterior

● -4.17

Variação face às últimas 6 semanas

● -1.92

Correções da semana (PP)

MAPE BIAS


Impacto nas previsões aprovadas da semana ● -0.70 ● -1.00

Impacto apenas nas previsões ajustadas ● -4.02 ● -10.71


Top correções positivas da semana (pp)



Variação semanal do MAPE RETEK aprovado (pp)



Variação semanal do BIAS RETEK aprovado (pp)



Gorgeous Ice Cream Case Study

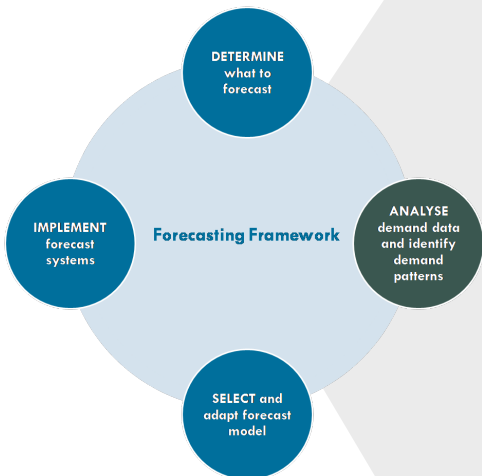


Outliers: atypical values caused by “rare” events or data entry errors


Analyze demand data and identify demand patterns

■ Outliers

<https://www.youtube.com/watch?v=UIAuPDpnwfy>




Before generating forecast, the historical baseline should be created



Many sectors have recently experienced dramatic changes on demand - relying solely on history is no longer valid

Demand Planning


Weekly sales



Prediction Intervals

We can create prediction intervals of the basic form:

“There is a xx% probability that the actual future value will fall in the range
Forecast \pm K (example: $K=1.96 \cdot \sigma(\text{forecast error})$, for 95% of confidence)





Carlson Department Store Case Study

- (1) the amount of sales Carlson would have made if the hurricane had not struck;
- (2) whether Carlson is entitled to any compensation for excess sales from increased business activity after the storm





Impact of aggregation level on forecasting performance



When comparing top-down with bottom-up forecasting processes it is often appropriate to measure forecasting accuracy both at aggregate and disaggregate level

Forecast Accuracy at different aggregation levels

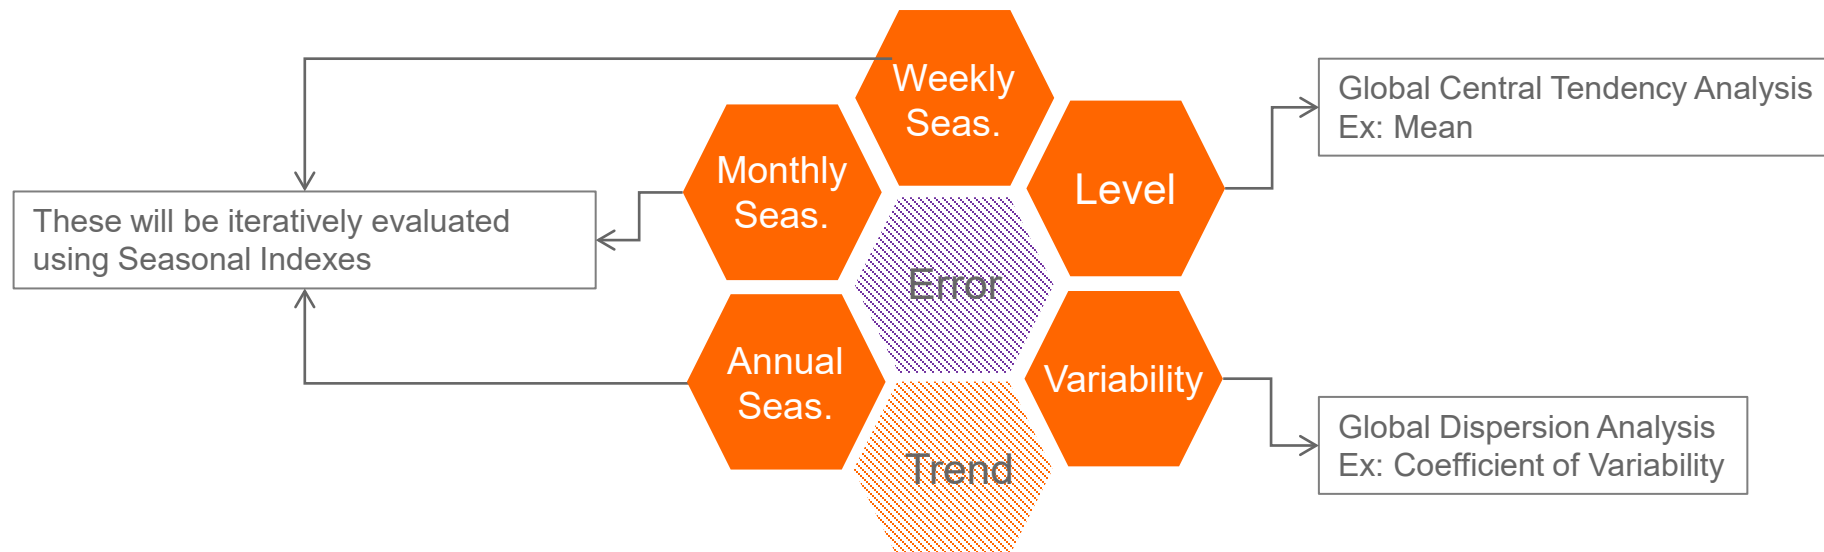


Impact of aggregation level on forecasting performance

Aggregation Policies

Clustering Process

- Select & Calculate Explanatory Variables – Item Clustering
 - The set of explanatory variables are the inputs for the clustering process to determine which observations are alike and which aren't;
 - Therefore, these inputs should be based on the time series components:



Impact of aggregation level on forecasting performance

Aggregation Policies


Clustering according to several likelihood variables.



- Honey & Sweets
- Deserts

864 SKU's

- ✓ CI.0 – 63 SKU's
- ✓ CI.1 – 133 SKU's
- ✓ CI.2 – 55 SKU's
- ✓ CI. 3 – 463 SKU's
- ✓ CI.4 – 150 SKU's



Impact of aggregation level on forecasting performance


Aggregation Policies

Clustering according to several likelihood variables.

- Water
- Beverages
- Soda

K-Means


- K=5
- Avg. Dist = -0.572
- Daves Bouldin = -1.277
- Example Dist. = 0.534



429 SKU's

- ✓ Cl.0 – 49 SKU's
- ✓ Cl.1 – 31 SKU's
- ✓ Cl.2 – 7 SKU's
- ✓ Cl.3 – 36 SKU's
- ✓ Cl.4 – 306 SKU's

Impact of aggregation level on forecasting performance



The Role of Time Series

Time Series Methods:
Exponential Smoothing

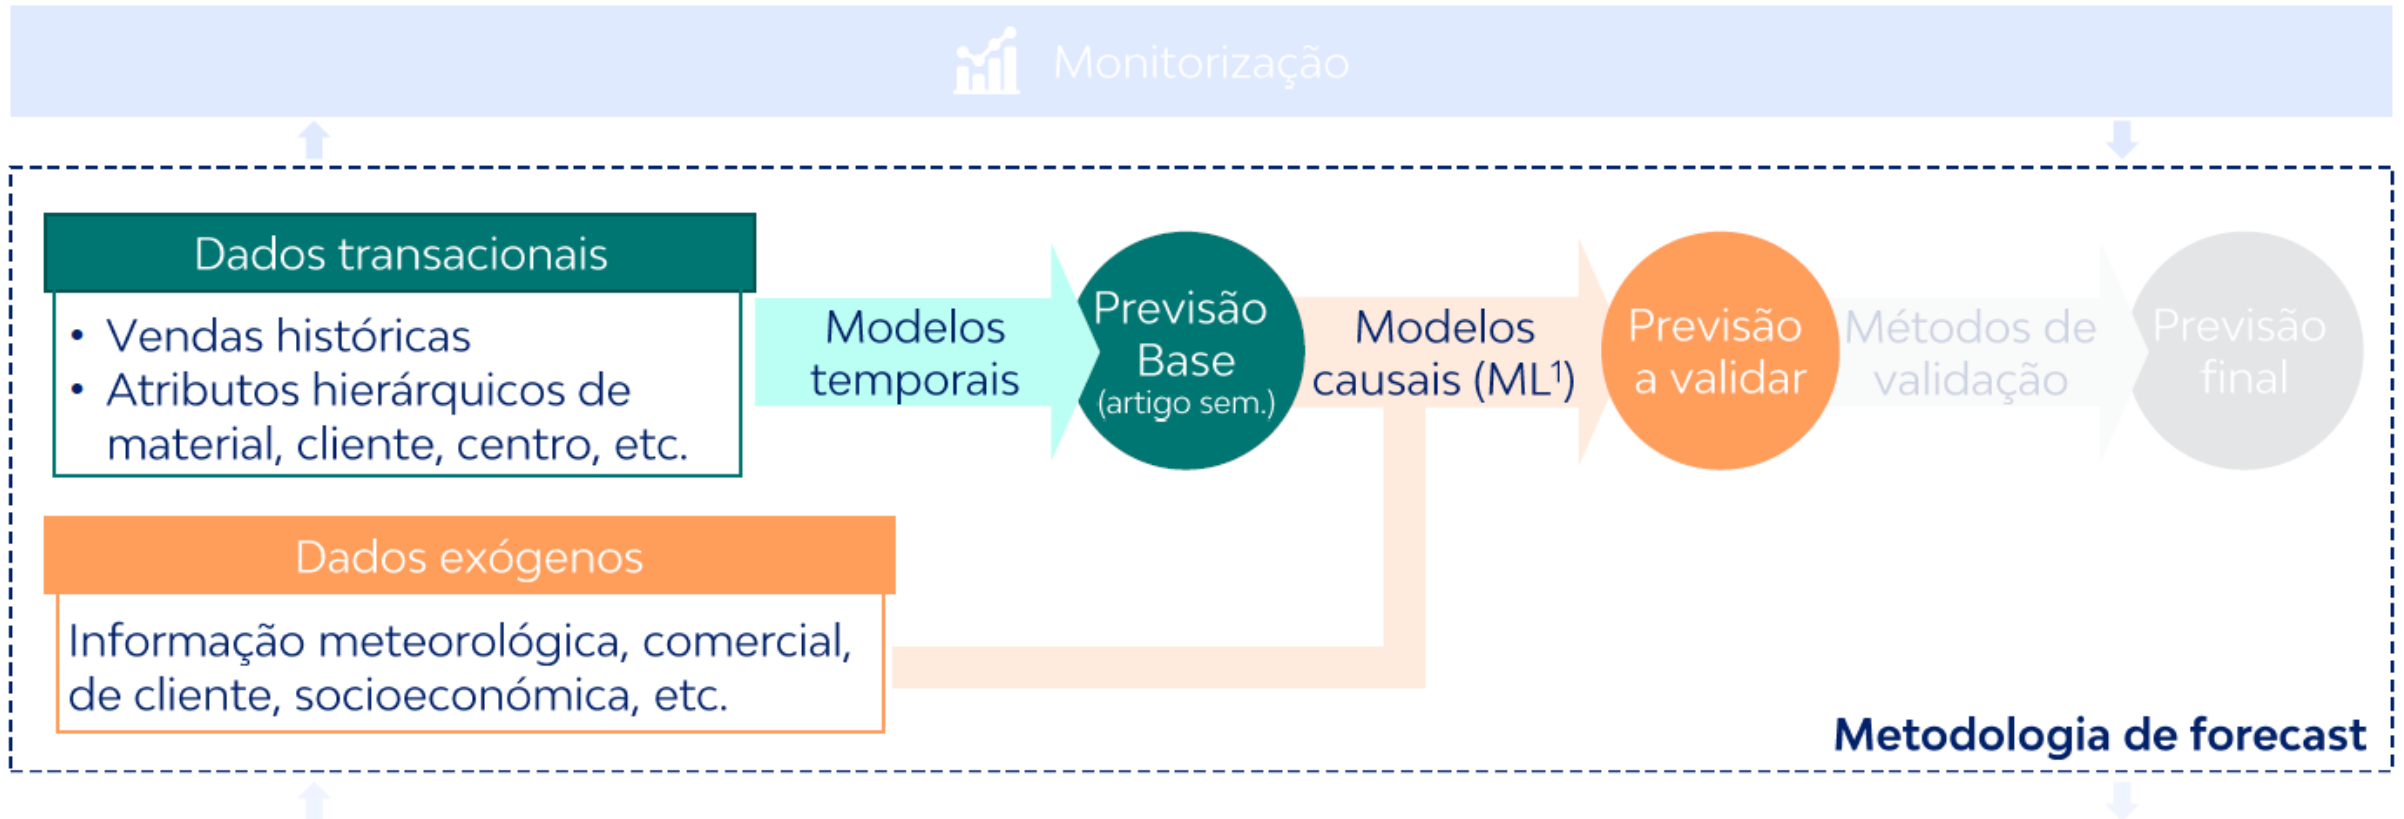
Accuracy, Outliers and
Aggregation

Hybrid Methods



Box & Jenkins



Using advanced analytic models with many sources of insights to understand how demand evolve



Using advanced analytic models with many sources of insights to understand how demand evolve

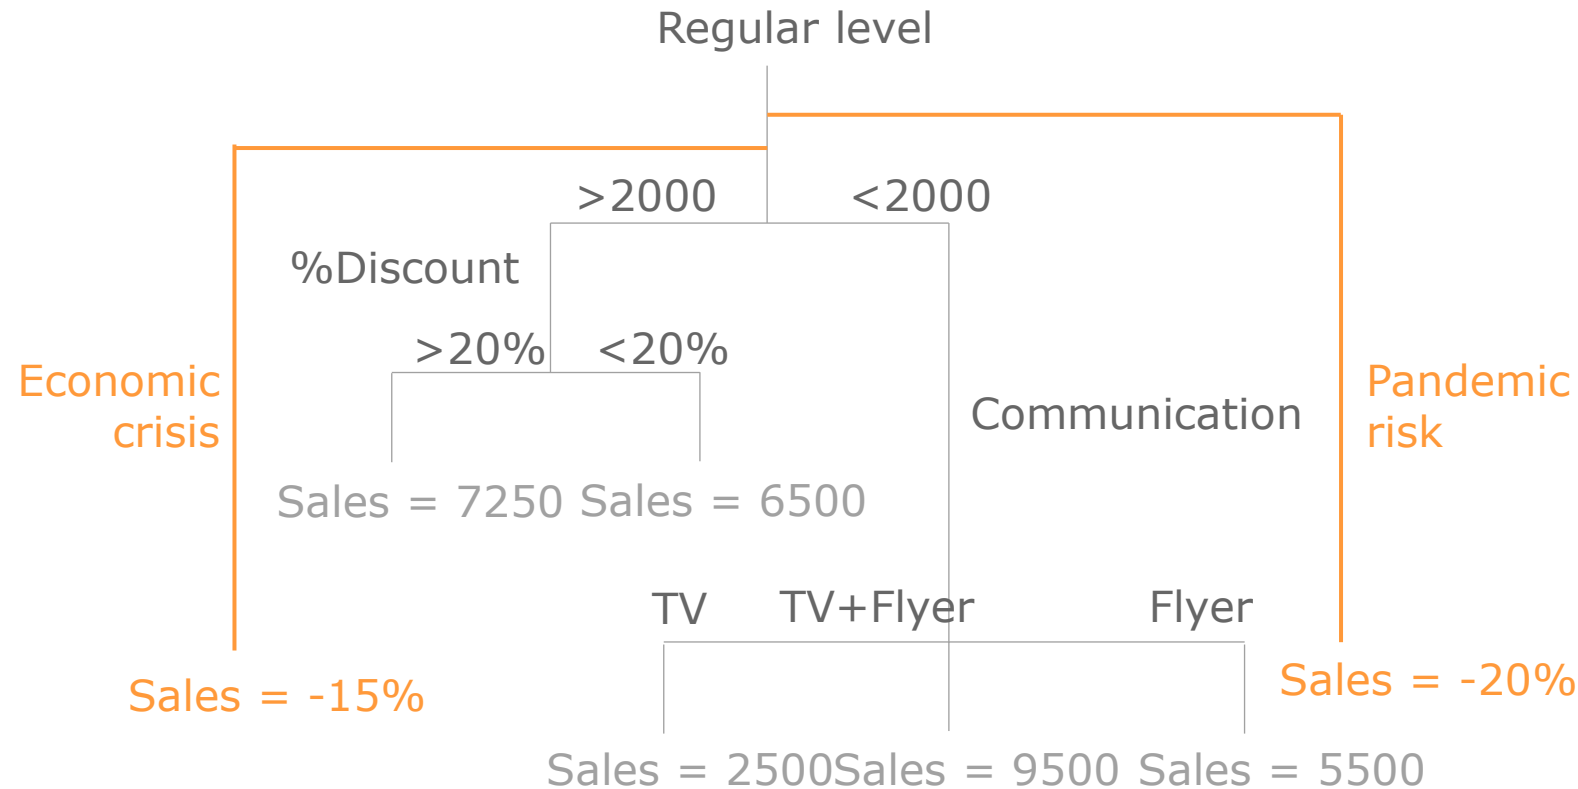
	Time-series	Machine learning (ML)		
	Holt-Winters + ensemble	GLM (Generalized linear model)	GBM (Gradient boosting machine)	RF ¹ (Random Forest)
 Vantagens	<ul style="list-style-type: none">• Capturam os efeitos temporais de uma série de vendas - nível, sazonalidade e tendência• Maior explicabilidade	<ul style="list-style-type: none">• Permitem a incorporação de variáveis exógenas para além da série de vendas, estimando a importância de cada uma para a previsão final• Self-learning dos parâmetros dos modelos com a alteração do histórico		
		Maior explicabilidade	Capturam efeitos não-lineares e interação entre variáveis	
 Desvantagens	<ul style="list-style-type: none">• Não permitem a incorporação de variáveis exógenas• Necessária atualização manual de parâmetros			
		Assumem a linearidade de efeitos	Podem levar a uma maior complexidade de interpretação	

Leading organizations are using advanced analytic models with many sources of insights to understand how demand evolve


ILLUSTRATIVE

Use advanced analytic models (e.g. Machine Learning) and **leverage internal and external sources** of information to **extract patterns and insights out of the past sales**

Price	Promotions	Weather
Consumer research	Social listening	Online Search Trends
Trends from other countries	Social economic scenario	



Using advanced analytic models with many sources of insights to understand how demand evolve



The Role of Time Series

Time Series Methods:
Exponential Smoothing

Accuracy, Outliers and
Aggregation

Hybrid Methods

Box & Jenkins



B&J: Base Idea of the method

WHATEVER THE FORECASTING MODELS, THE FORECASTS ARE OBTAINED FROM THE AVAILABLE OBSERVATIONS, ACCORDING TO A FUNCTION OF TYPE

$$\hat{Z}_t(k) = f(Z_t, Z_{t-1}, Z_{t-2}, \dots)$$

IN THE PREVIOUS METHODS, EACH OF THEM BASED ON ITS OWN, THE ADOPTION OF A METHOD INVOLVED THE ESTABLISHMENT OF THE BASIC STRUCTURE OF THE RELATIONSHIP BETWEEN FORECASTS AND AVAILABLE DATA

$$\hat{Z}_t(1) = \alpha \cdot Z_t + \alpha \cdot (1 - \alpha) \cdot Z_{t-1} + \dots$$

FOR EXAMPLE, IN THE SIMPLE EXPONENTIAL SMOOTHING METHOD ONLY ONE DEGREE OF FREEDOM WAS AVAILABLE TO ADAPT THE MODEL TO REALITY: THE CHOICE OF α

B&J: Base Idea of the method

IN THE APPROACH PROPOSED BY BOX AND JENKINS

- A WIDE RANGE OF ALTERNATIVE MODELS IS CONSIDERED
- ON THE BASIS OF HISTORICAL OBSERVATIONS, IT IS CALCULATED THAT DIFFERENT STATISTICS ARE CARRIED OUT ON THE BASIS OF THESE DIFFERENT TESTS THAT ALLOW YOU TO CHECK WHICH OF THE MODELS IT IS THE ONE THAT BEST ADAPTS TO THE AVAILABLE DATA


B&J Models applicable to stationary time series

Autoregressive Models


- THE SIMPLEST MODEL THAT TRANSLATES THE EXISTENCE OF CORRELATION BETWEEN THE SUCCESSIVE TERMS OF A SERIES IS

$$Z_t = \phi_1 \cdot Z_{t-1} + E_t \quad \text{COM } E_t \sim \text{IN}(0, \sigma_E^2)$$

$$\phi_1 > 0$$



$$\phi_1 < 0$$



B&J Models applicable to stationary time series

Autoregressive Models

/...

ANOTHER, MORE COMPLICATED MODEL FOR TRANSLATING THE EXISTENCE OF CORRELATION BETWEEN THE SUCCESSIVE TERMS OF A SERIES WOULD BE

$$Z_t = \phi_1 \cdot Z_{t-1} + \phi_2 \cdot Z_{t-2} + E_t \quad \text{COM } E_t \sim \text{IN}(0, \sigma_E^2)$$

IN GENERAL, AN AUTOREGRESSIVE MODEL OF ORDER p CAN BE DEFINED AS FOLLOWS:

$$\text{AR}(p): \quad Z_t = \phi_1 \cdot Z_{t-1} + \phi_2 \cdot Z_{t-2} + \dots + \phi_p \cdot Z_{t-p} + E_t \quad \text{COM } E_t \sim \text{IN}(0, \sigma_E^2)$$

THE DESIGNATION OF THE MODEL - AUTOREGRESSIVE - HAS ITS ORIGIN IN THE FACT THAT IT IS A REGRESSION OF THE VARIABLE Z_t IN FUNCTION OF ITSELF, BUT WITH A LAG 1, 2, ... , p .../

B&J Models applicable to stationary time series

Autoregressive Models

A MORE SUCCINCT WAY OF REPRESENTING AN AR(p) MODEL IS ACHIEVED BY USING THE 'BACK SHIFT OPERATOR'

IT IS AN OPERATOR THAT, WHEN APPLIED TO THE REALIZATION OF ANY VARIABLE AT THE TIME t TRANSFORMS IT INTO THE REALIZATION OF THAT VARIABLE IN $t-1$

$$B Z_t = Z_{t-1}$$

$$B^2 Z_t = B(B Z_t) = B Z_{t-1} = Z_{t-2}$$

(...)

$$B^j Z_t = Z_{t-j}$$

B&J Models applicable to stationary time series

Autoregressive Models

/...

USING OPERATOR B , THE $AR(p)$ TURNS INTO

$$Z_t = \phi_1 Z_{t-1} + \phi_2 Z_{t-2} + \dots + \phi_p Z_{t-p} + E_t$$

$$Z_t - \phi_1 Z_{t-1} - \phi_2 Z_{t-2} - \dots - \phi_p Z_{t-p} = E_t$$

$$(1 - \phi_1 B - \phi_2 B^2 - \dots - \phi_p B^p) Z_t = E_t$$

$$AR(p): \phi(B) Z_t = E_t$$

WHERE $\phi(B)$ REPRESENTS A POLYNOMIAL OF DEGREE p IN B

B&J Models applicable to stationary time series

Moving Average Models

AUTOREGRESSIVE MODELS CAN BE MANIPULATED AND EXPRESSED THE VARIABLE Z_t AS A FUNCTION OF ERRORS RECORDED AT TIME t OR AT EARLIER INSTANTS

EXAMPLE

$$\begin{aligned}\text{AR}(1): \quad Z_t &= \phi_1 Z_{t-1} + E_t \\ &= \phi_1 \cdot (\phi_1 Z_{t-2} + E_{t-1}) + E_t \\ &= \phi_1^2 Z_{t-2} + \phi_1 E_{t-1} + E_t \\ &= \phi_1^2 \cdot (\phi_1 Z_{t-3} + E_{t-2}) + \phi_1 E_{t-1} + E_t \\ &= E_t + \phi_1 E_{t-1} + \phi_1^2 E_{t-2} + \phi_1^3 E_{t-3} + \dots\end{aligned}$$

B&J Models applicable to stationary time series

Moving Average Models

- .../

IN A MOVING AVERAGE MODEL OF ORDER q , THE VARIABLE z_t IS EXPRESSED AS A FUNCTION OF THE ERRORS RECORDED AT THE TIME t AND IN THE PREVIOUS MOMENTS

A MODEL OF THE MOVING AVERAGE TYPE OF ORDER q IS THEN GIVEN BY

$$\begin{aligned}\text{MA}(q): \quad Z_t &= E_t + \theta_1 \cdot E_{t-1} + \theta_2 \cdot E_{t-2} + \dots + \theta_q \cdot E_{t-q} \\ &= (1 + \theta_1 \cdot B + \theta_2 \cdot B^2 + \dots + \theta_q \cdot B^q) \cdot E_t \\ &= \theta(B) \cdot E_t\end{aligned}$$

WHERE $q(B)$ REPRESENTS A POLYNOMIAL OF DEGREE q IN B

AN AR MODEL WITH A FINITE NUMBER OF TERMS (AND THEREFORE WITH A FINITE NUMBER OF PARAMETERS) IS EQUIVALENT TO AN MA MODEL WITH AN INFINITE NUMBER OF PARAMETERS

B&J Models applicable to stationary time series

Mixed Models

MIXED MODELS ARE MODELS THAT SIMULTANEOUSLY HAVE AN AR(P) COMPONENT AND AN MA(q) COMPONENT, BEING DENOTED MODELS(p,q)

$$\text{ARMA}(p,q): \phi(B) \cdot Z_t = \theta(B) \cdot E_t$$

$$(1 - \phi_1 \cdot B - \phi_2 \cdot B^2 - \dots - \phi_p \cdot B^p) \cdot Z_t = (1 + \theta_1 \cdot B + \theta_2 \cdot B^2 + \dots + \theta_q \cdot B^q) \cdot E_t$$

ARMA(p,q) MODELS (WITH A FINITE NUMBER OF PARAMETERS) ARE ALSO CONVERTIBLE INTO AR OR MA MODELS WITH INFINITE NUMBER OF PARAMETERS

B&J Models applicable to stationary time series

Model Selection

THE AIM OF THIS PHASE IS, ON THE ONE HAND, TO IDENTIFY, AMONG THE BJ MODELS, WHICH ONE(S) BEST SUITS EACH SERIES AND, ON THE OTHER HAND, TO PRELIMINARILY ESTIMATE THE PARAMETERS OF THE IDENTIFIED MODEL(S)

THE IDENTIFICATION OF THE MODEL(S) IS CARRIED OUT ON THE BASIS OF THE BEHAVIOR OF THE AUTOCORRELATION AND PARTIAL AUTOCORRELATION

$$\hat{\phi}_{kk} \sim \mathbf{N}\left(\mu = 0, \sigma^2 \approx \frac{1}{N}\right), \text{ PARA } k > p \quad \Rightarrow \quad \text{MODELO AR}(p)$$

$$r_k \sim \mathbf{N}\left[\mu = 0, \sigma^2 \approx \frac{1}{N} \cdot \left(1 + 2 \cdot \sum_{i=1}^q r_i^2\right)\right], \text{ PARA } k > q \quad \Rightarrow \quad \text{MODELO MA}(q)$$

$$\text{NEM } \hat{\phi}_{kk} \text{ NEM } r_k \text{ SOFREM UM «CUT OFF»} \quad \Rightarrow \quad \text{MODELO ARMA}$$


B&J Models applicable to stationary time series

• q .../

Autocorrelation Function (f.a.c.)

IN GENERAL, FOR A MA(q) MODEL, THE AUTOCORRELATIONS CANCEL EACH OTHER OUT FOR $k > q$; KNOWN AS 'CUT OFF', INSTEAD OF EXPONENTIAL DECAY

AS AN EXAMPLE, FOR A MODEL MA(2), THE PATTERNS THAT THE AUTOCORRELATION FUNCTION CAN DISPLAY ARE PRESENTED




B&J Models applicable to stationary time series

Partial Autocorrelation Function (f.a.c.p.)

IN GENERAL, FOR A $AR(p)$ MODEL, THE PARTIAL AUTOCORRELATIONS CANCEL EACH OTHER OUT FOR $k > p$; KNOWN AS 'CUT OFF', INSTEAD OF EXPONENTIAL DECAY

AS AN EXAMPLE, FOR A MODEL $AR(2)$, THE PATTERNS THAT THE PARTIAL AUTOCORRELATION FUNCTION CAN DISPLAY ARE PRESENTED



B&J Models applicable to stationary time series

Model Selection

HYPOTHESIS TESTING OF AUTOCORRELATIONS

SUCCESSIVELY FOR $q = 0, 1, 2, \dots$, CHECK IF $\rho_k = 0$ PARA $k > q$

$$\begin{aligned} H_0: \rho_k &= 0 \\ H_1: \rho_k &\neq 0 \\ ET &= \frac{r_k}{\sqrt{\frac{1}{N} \cdot \left(1 + 2 \cdot \sum_{i=1}^q r_i^2 \right)}} \\ H_0 \text{ VERDADEIRA} &\Rightarrow ET \sim N(0,1) \end{aligned}$$

BASED ON THIS TEST, IDENTIFY WHICH AUTOCORRELATIONS ARE SIGNIFICANTLY DIFFERENT FROM ZERO

B&J Models applicable to stationary time series

Model Selection

HYPOTHESIS TESTING OF PARTIAL AUTOCORRELATIONS

VERIFY IF $\phi_{kk} = 0$ PARA $k > p$

$$H_0: \phi_{kk} = 0$$

$$H_1: \phi_{kk} \neq 0$$

$$ET = \frac{\hat{\phi}_{kk}}{\sqrt{\frac{1}{N}}}$$



$$H_0 \text{ VERDADEIRA} \Rightarrow ET \sim N(0, 1)$$

BASED ON THIS TEST, IDENTIFY WHICH PARTIAL AUTOCORRELATIONS ARE SIGNIFICANTLY DIFFERENT FROM ZERO

B&J Models applicable to non-stationary time series

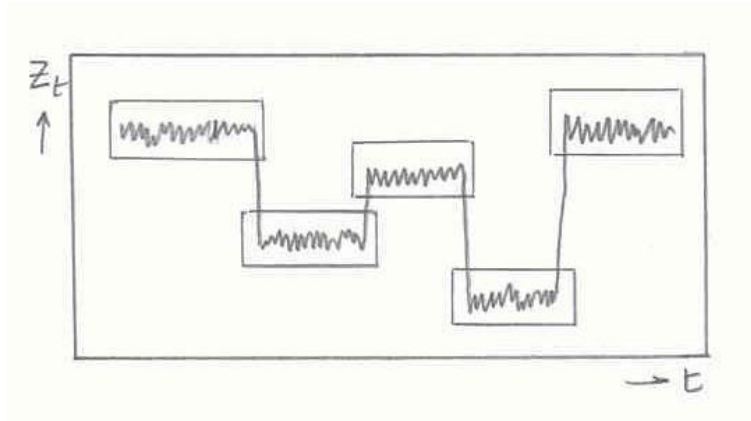
Non-stationary signals

In case the non-stationarity origins from level and trend variations or seasonalities, it is revealed from the f.a.c. as it does not present cut-offs for low lag values.



B&J Models applicable to non-stationary time series

Local stationary time series with "jumps"



Locally


$$E(Z_t) \approx \text{CONSTANTE}$$

$$\Rightarrow E(Z_t - Z_{t-1}) \approx 0$$

The differentiation of the original time series transforms it in a stationary time series

$$Y_t = Z_t - Z_{t-1} = \nabla Z_t$$

$$E(Y_t) \approx 0$$



$$Y_t = \phi_1 \cdot Y_{t-1} + E_t$$

$$Z_t - Z_{t-1} = \phi_1 \cdot (Z_{t-1} - Z_{t-2}) + E_t$$


$$\text{AR}(1) \quad Z_t = (1 + \phi_1) \cdot Z_{t-1} - \phi_1 \cdot Z_{t-2} + E_t$$

$$Z_{t+1} = (1 + \phi_1) \cdot Z_t - \phi_1 \cdot Z_{t-1} + E_{t+1}$$

$$\hat{Z}_t(1) = (1 + \phi_1) \cdot z_t - \phi_1 \cdot z_{t-1} + E_{t+1} = (1 + \phi_1) \cdot z_t - \phi_1 \cdot z_{t-1}$$

B&J Models applicable to non-stationary time series

Time series with local linear trends



Locally

$$E(Z_t) \approx \alpha_0 + \alpha_1 \cdot t$$

$$\Rightarrow E(Z_t - Z_{t-1}) \approx \alpha_1$$

$$\Rightarrow E[(Z_t - Z_{t-1}) - (Z_{t-1} - Z_{t-2})] \approx 0$$


Double differentiation

$$Y_t = \nabla^2 Z_t = \nabla(Z_t - Z_{t-1}) = (Z_t - Z_{t-1}) - (Z_{t-1} - Z_{t-2}) = Z_t - 2 \cdot Z_{t-1} + Z_{t-2}$$

$$E(Y_t) \approx 0$$


B&J Models applicable to non-stationary time series

Seasonal time series: Seasonal Autoregressive Integrated Moving Average



$$E(Z_t) \approx \alpha_0 + f_t \quad (f_t: \text{COMP. SAZONAL ADITIVA, COM PERIODO } s)$$

$$\Rightarrow E(Z_t - Z_{t-s}) \approx 0 \quad Y_t = \nabla_s Z_t = Z_t - Z_{t-s}$$



$$E(Z_t) \approx \alpha_0 + \alpha_1 \cdot t + f_t$$

$$\Rightarrow E(Z_t - Z_{t-s}) \approx s \cdot \alpha_1 \quad Y_t = \nabla \nabla_s Z_t = (Z_t - Z_{t-s}) - (Z_{t-1} - Z_{t-s-1})$$

$$\Rightarrow E[(Z_t - Z_{t-s}) - (Z_{t-1} - Z_{t-s-1})] \approx 0$$

$$\text{SARIMA}(p, d, q) \times (P, D, Q): \Phi_P(B^T) \cdot \phi_p(B) \cdot \nabla^d \cdot \nabla_s^D Z_t = \Theta_Q(B^T) \cdot \theta_q(B) \cdot E_t$$

(S)ARIMA for Python

```
1 from pandas import read_csv
2 from pandas import datetime
3 from matplotlib import pyplot
4 from pandas.plotting import autocorrelation_plot
5
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8
9 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
10 autocorrelation_plot(series)
11 pyplot.show()
```

statsmodels.tsa.arima.model.ARIMA

https://alkaline-ml.com/pmdarima/index.html

ndarima
0.3

» pmdarima: ARIMA estimators for Python [View page source](#)

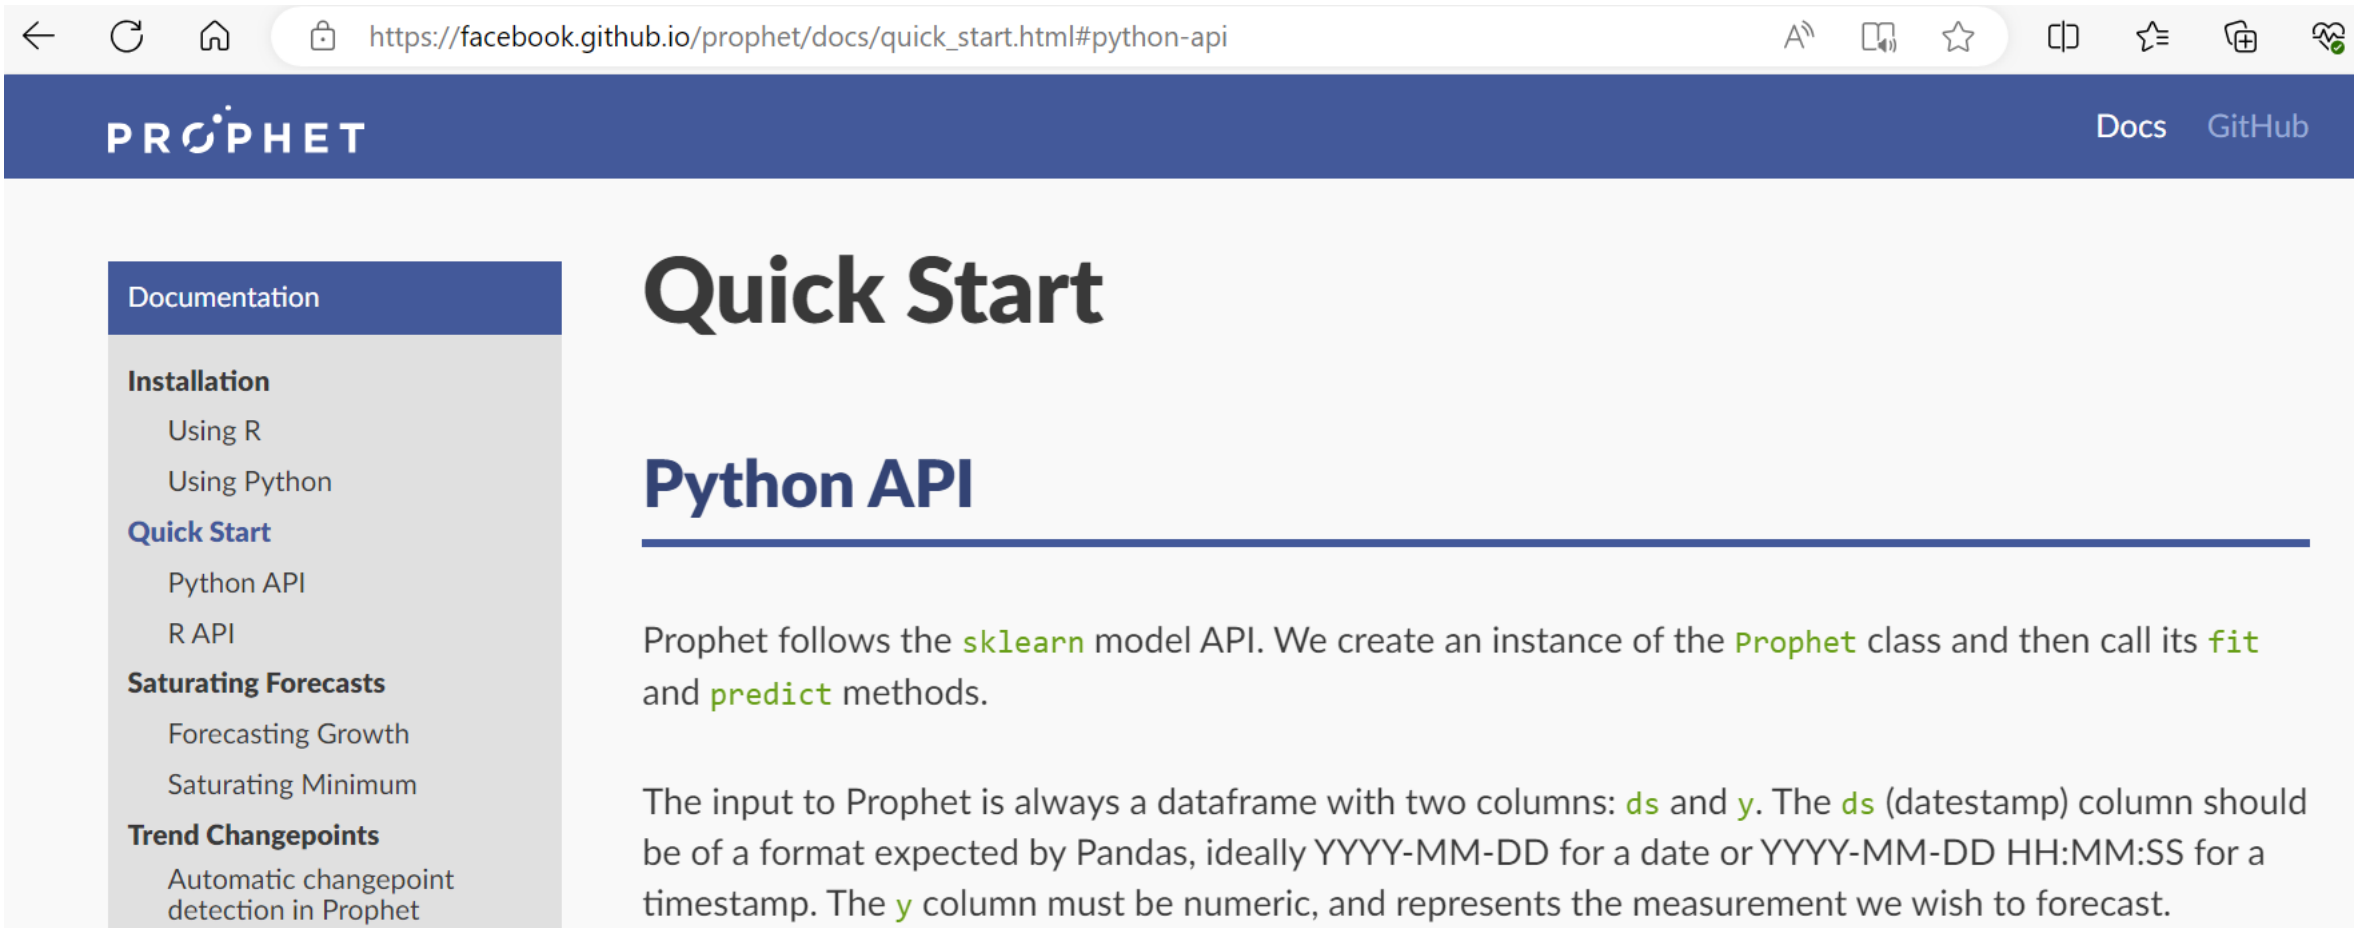
pmdarima: ARIMA estimators for Python

`pmdarima` brings R's beloved `auto.arima` to Python, making an even stronger case for why you don't need R for data science. `pmdarima` is 100% Python + Cython and does not leverage any R code, but is implemented in a powerful, yet easy-to-use set of functions & classes that will be familiar to scikit-learn users.

```
class statsmodels.tsa.arima.model.ARIMA(
    endog,
    exog=None,
    order=(0, 0, 0),
    seasonal_order=(0, 0, 0, 0),
    trend=None,
    enforce_stationarity=True,
    enforce_invertibility=True,
    concentrate_scale=False,
    trend_offset=1,
    dates=None,
    freq=None,
    missing='none',
    validate_specification=True
)
```

[\[source\]](#)

Prophet - procedure for forecasting time series data based on an additive model where non-linear trends are fit with yearly, weekly, and daily seasonality,



The screenshot shows a web browser displaying the Prophet Python API documentation. The browser's address bar shows the URL: https://facebook.github.io/prophet/docs/quick_start.html#python-api. The page has a dark blue header with the 'PROPHET' logo on the left and 'Docs' and 'GitHub' links on the right. A left sidebar contains a 'Documentation' menu with items: 'Installation' (sub-items: 'Using R', 'Using Python'), 'Quick Start' (sub-items: 'Python API', 'R API'), 'Saturating Forecasts' (sub-items: 'Forecasting Growth', 'Saturating Minimum'), and 'Trend Changepoints' (sub-item: 'Automatic changepoint detection in Prophet'). The main content area has a 'Quick Start' section followed by a 'Python API' sub-section. The text under 'Python API' states: 'Prophet follows the `sklearn` model API. We create an instance of the `Prophet` class and then call its `fit` and `predict` methods.' Below this, it explains the input: 'The input to Prophet is always a dataframe with two columns: `ds` and `y`. The `ds` (datestamp) column should be of a format expected by Pandas, ideally YYYY-MM-DD for a date or YYYY-MM-DD HH:MM:SS for a timestamp. The `y` column must be numeric, and represents the measurement we wish to forecast.'

Make change happen